

# Machine Learning for IoT

Part-I: Technologies for the IoT  
(HW perspective)

---

# Objectives

---

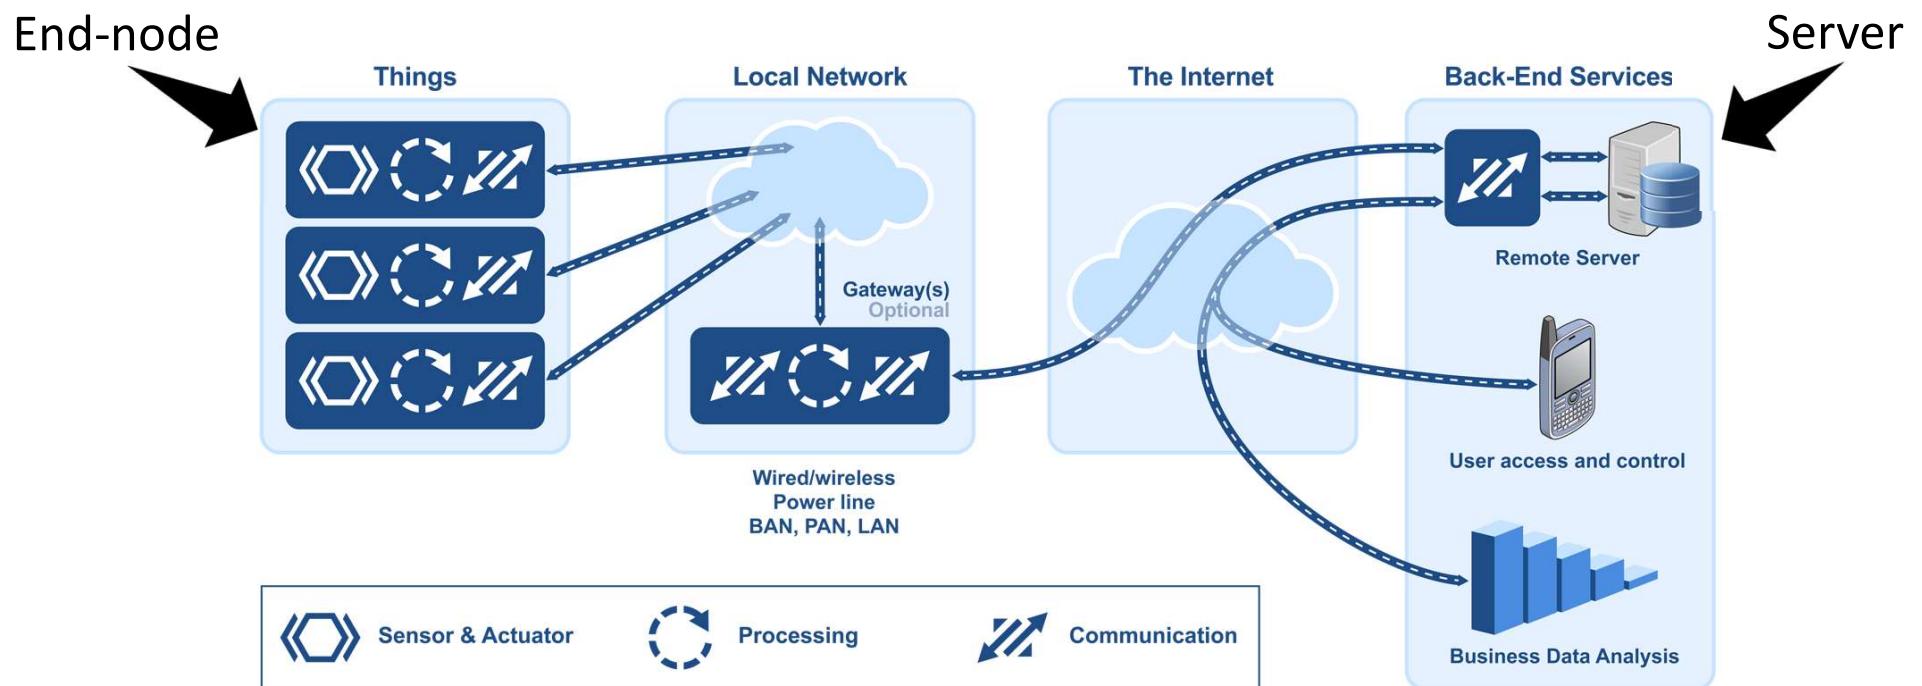
- An overview of the HW devices in a typical IoT environment
- Taxonomy of HW components and their functionality
  - The ***functionality*** you can accommodate in your system
  - The computational power you need for a given ***functionality***
  - The ***power/energy*** sizing of your system
- Specifically, assuming you are not a “HW expert”

# Contents

---

- Architectural overview
- Sensors (and actuators)
- Signal conditioning
- Processing units: MCU vs MPU
- Radio communication
- Beyond functionality: Energy path

# IoT infrastructure

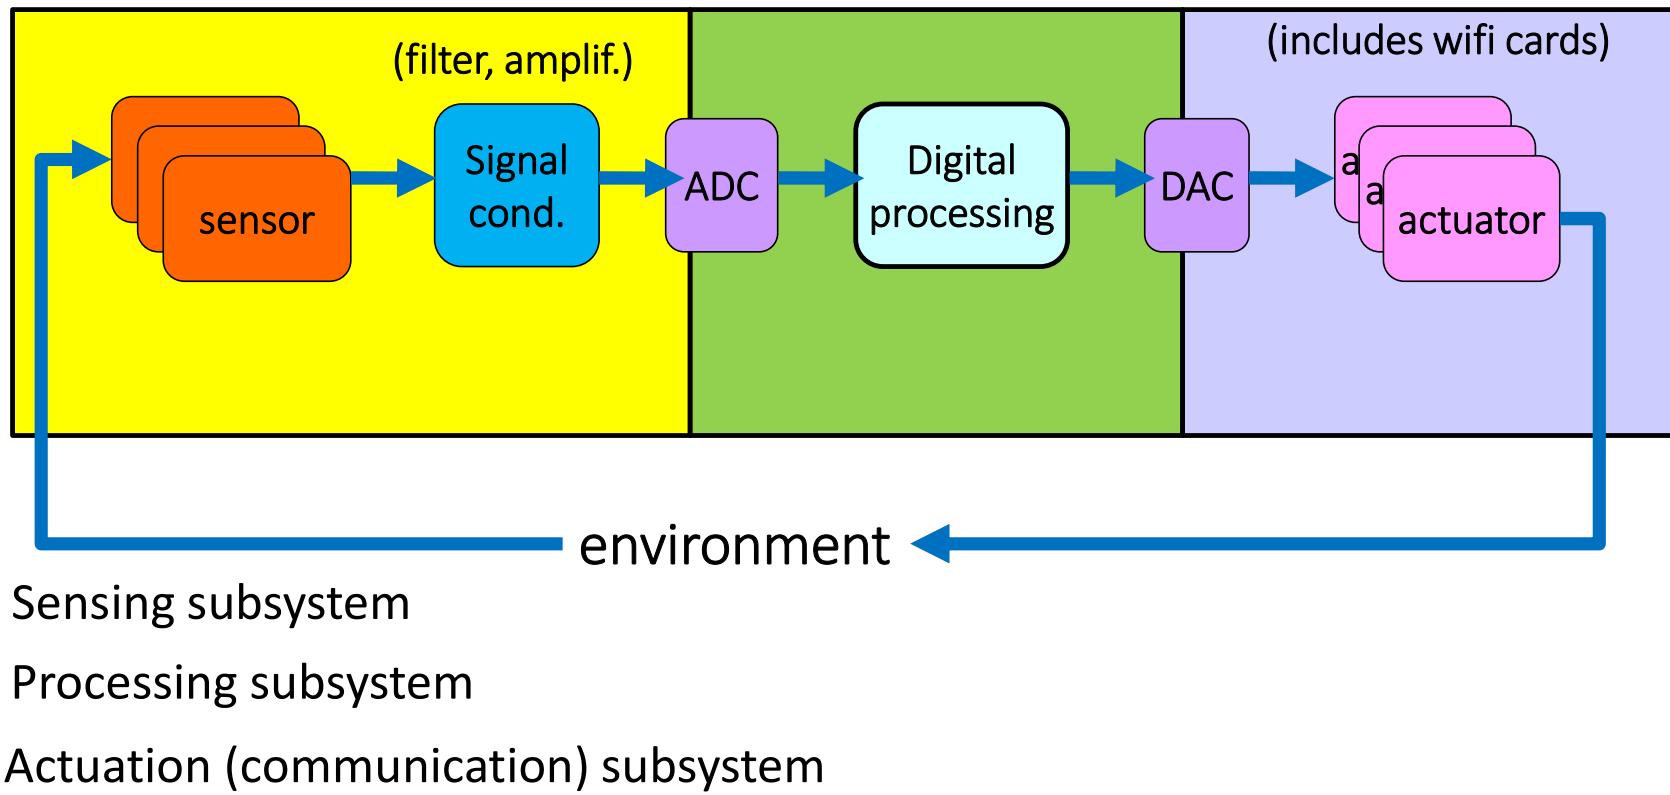


# End-nodes vs server/cloud nodes

---

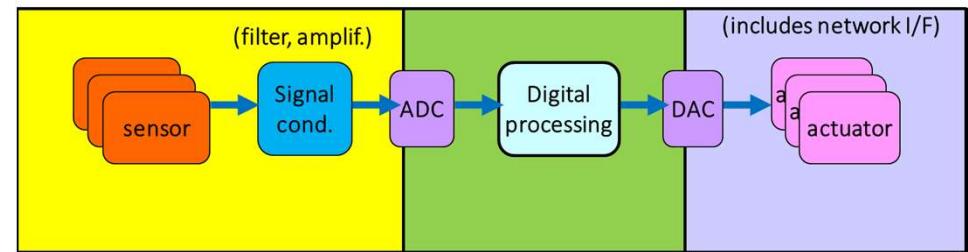
- Form Factor
  - End-nodes:  $\text{cm}^2/\text{mm}^2$
  - Server/Cloud:  $(10^2)\text{m}^2$
- Functionality
  - End-nodes: Cyber-physical-system, with sensing, processing, acting capability, wifi connectivity
  - Server/cloud: (mostly) processing and storage, wired connectivity
- Performance
  - End-nodes: MHz/MFLOPs, MIPs
  - Server/cloud: GHz/many GFLOPs to TFLOPs
- Power/Energy budget
  - End-nodes: mw/W ( $<<7\text{W}$ )
  - Server/cloud: kW/MW

# Functional view



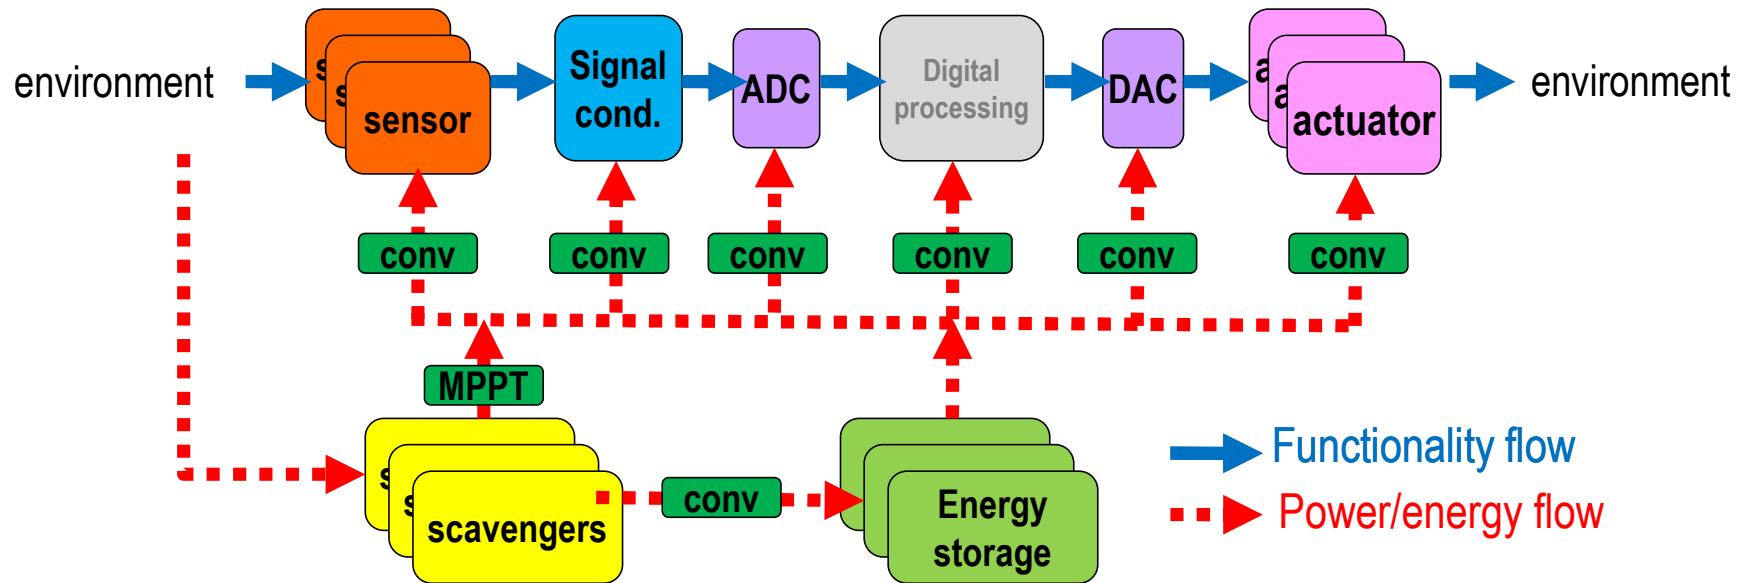
# (Extra) Functionality

- This is a functional view
  - What the system does in terms of functionality
  - Connections represent flow of data



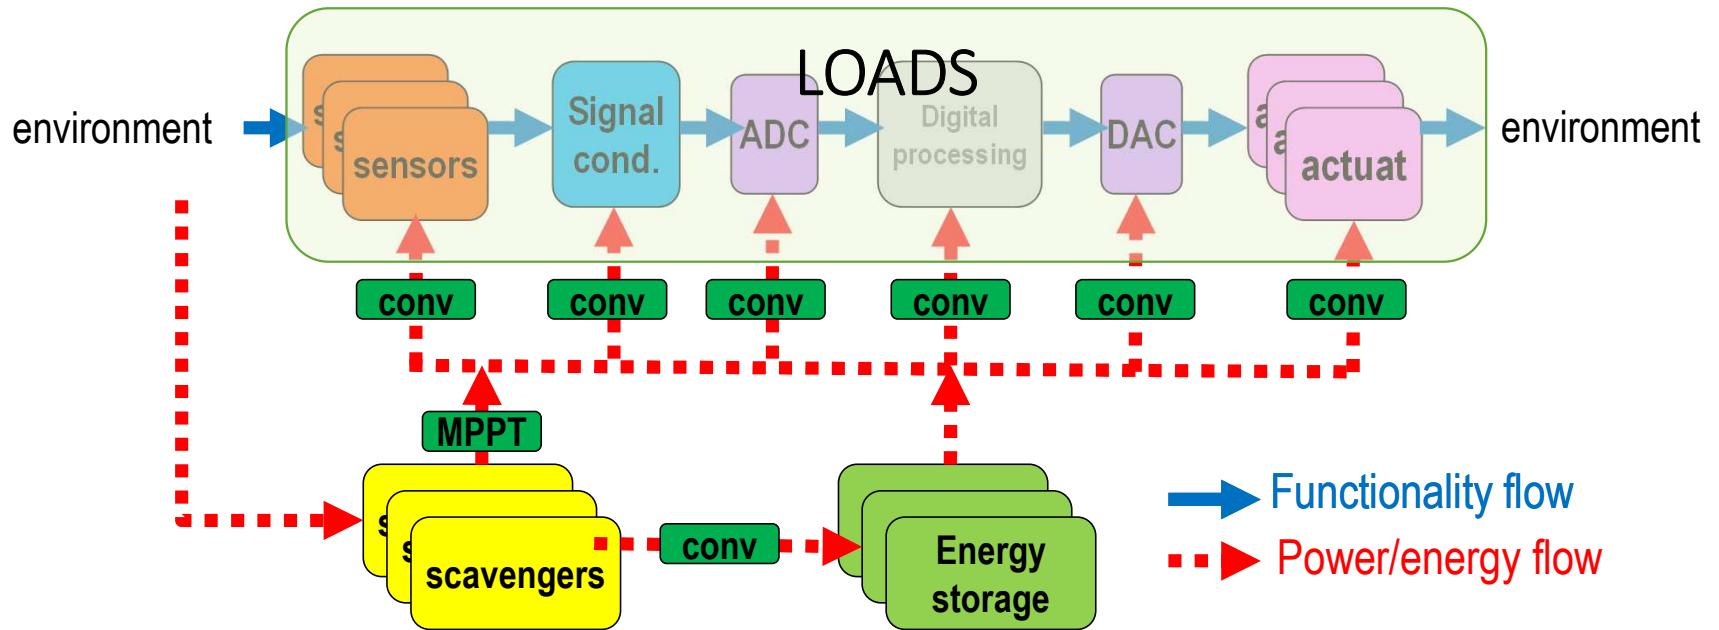
- What is missing?
- A Non-functional view
  - Most important non-functional information is the flow of energy/power
  - Implies adding other components

# Functional view: info vs energy/power flows

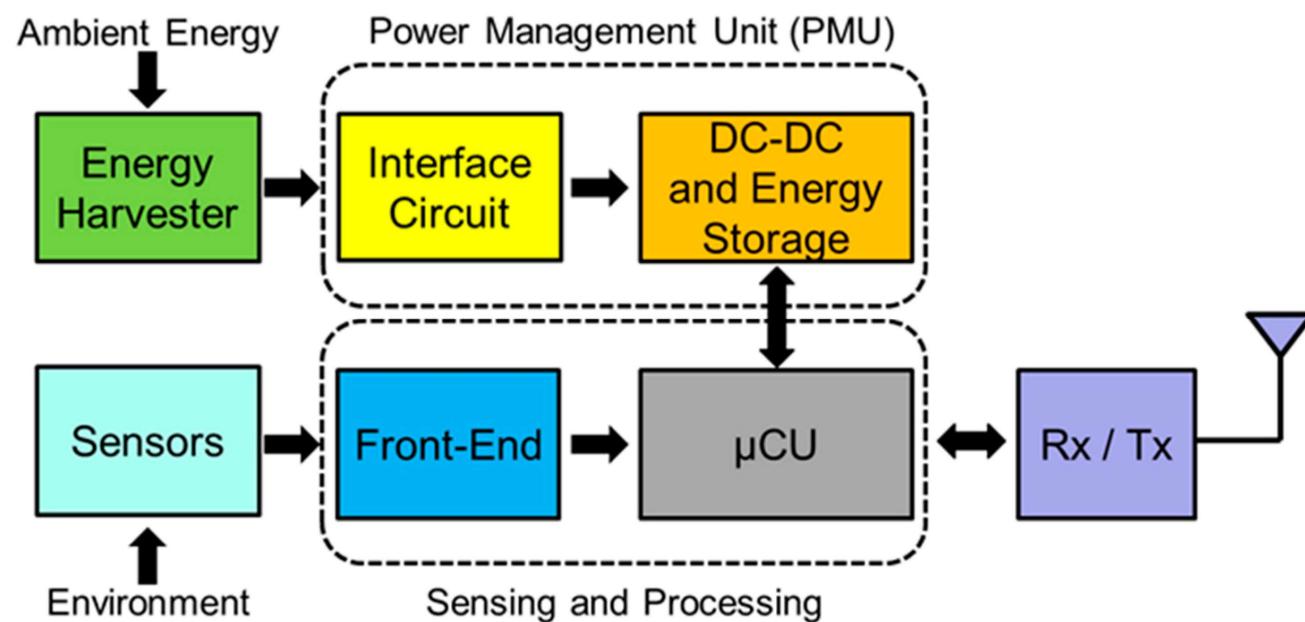


# Functional view: info vs energy/power flows

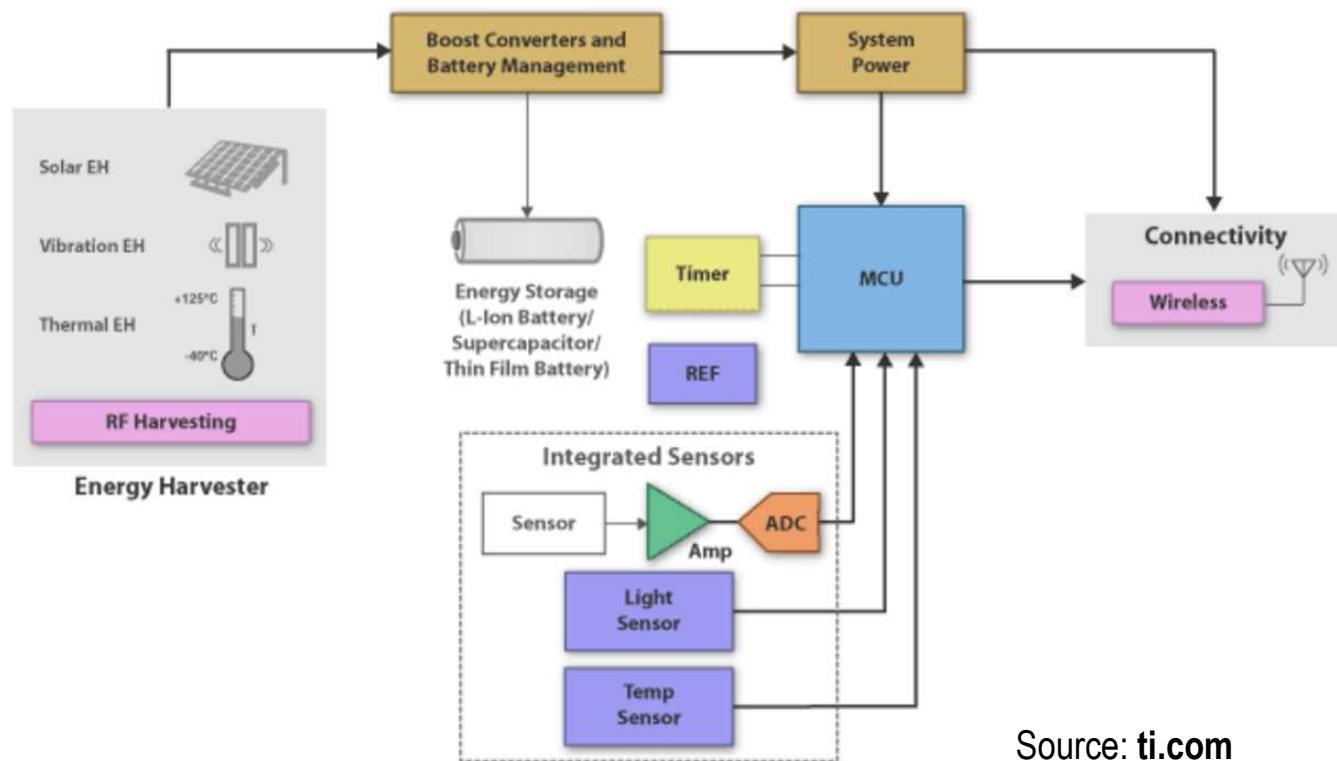
- Functional components are generic ‘loads’ from the power flow perspective!!!



# Structural view



# Integrated solutions



Source: [ti.com](http://ti.com)

# Lecture plan

---

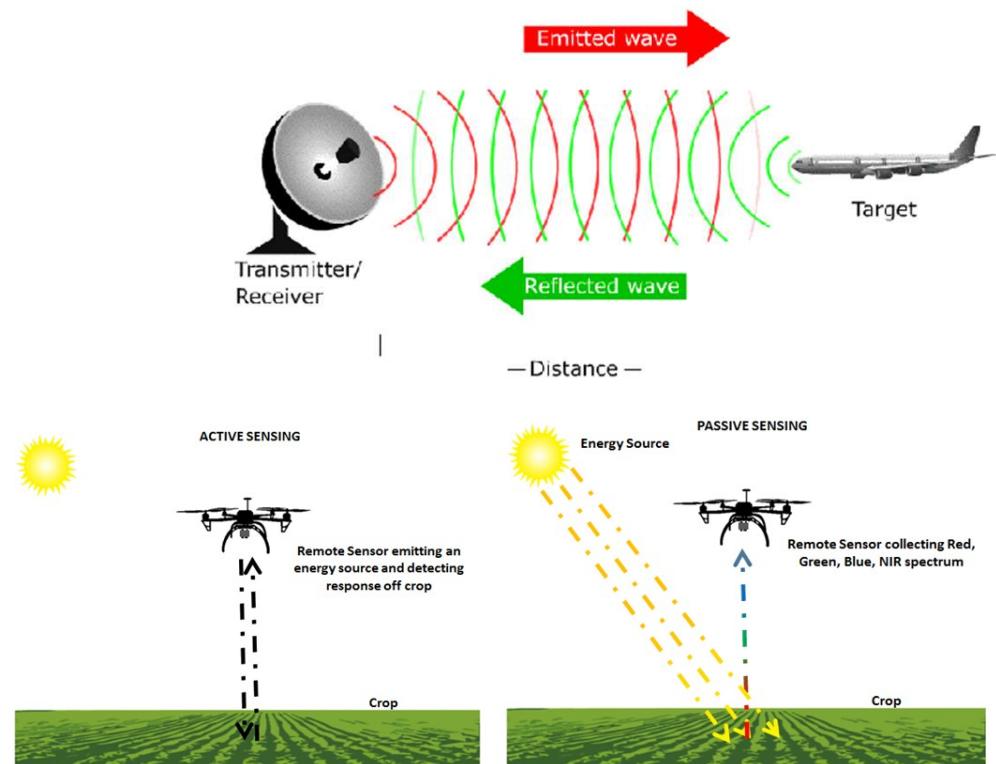
- First explore the functional domain
  - Sensing subsystem
  - Processing subsystem
  - Communication subsystem
- Then the extra-functional (power) domain
  - Understand the issues involving energy
    - Consuming (as little as possible) energy
    - Generating energy
    - Storing energy
    - Convert energy
- Dive into computers architectures (as separate section)
  - Understand how commercial cores evolved over time
  - Taxonomy of existing programming/executions models and their characteristics
  - What are the options available today for complex ML-powered IoT

## Sensors and acquisition

---

# Active vs passive sensors

- Active = require an external power supply to operate, called an excitation signal which is used by the sensor to produce the output signal.
  - Example:
    - Radar
    - Lidar
- Passive = does not need any additional power source or excitation voltage.
  - They generates an output signal in response to some external stimulus
  - Example:
    - a thermocouple which generates its own voltage output when exposed to heat
    - Image sensor (camera)



<https://towardsdatascience.com/why-tesla-wont-use-lidar-57c325ae2ed5>

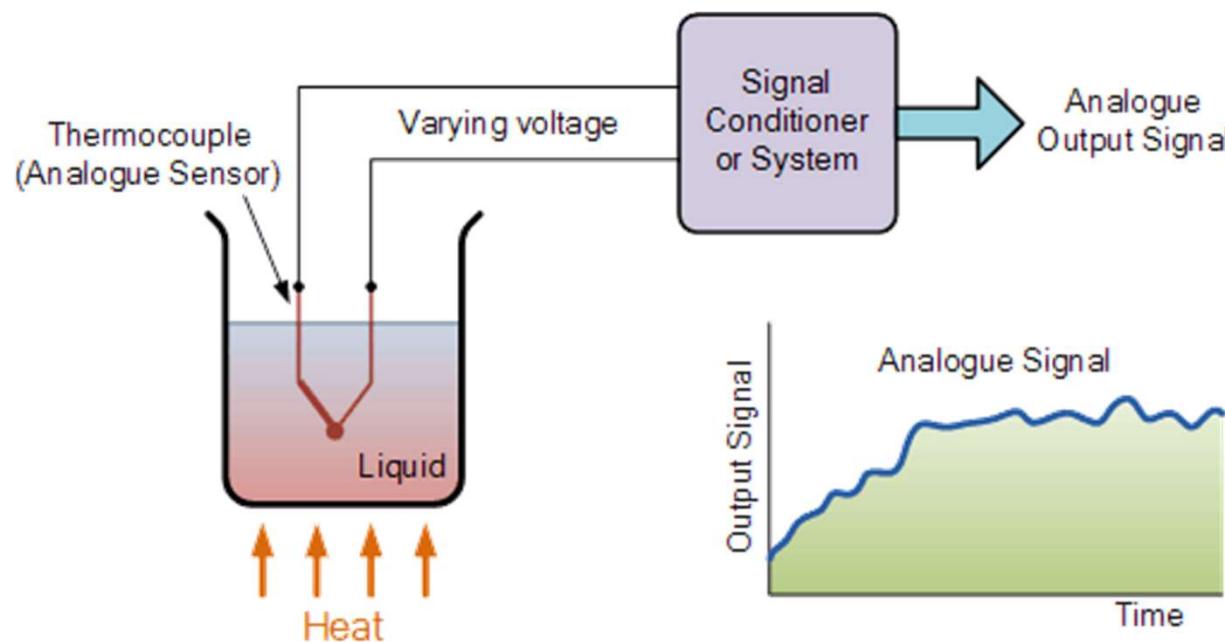
# Analog vs Digital sensors

---

- Analog: produce a continuous output signal or voltage which is generally proportional to the quantity being measured (most common)
- Digital (or discrete): produce a discrete digital output signals or voltages that are a digital representation of the quantity being measured.

# Analog sensor

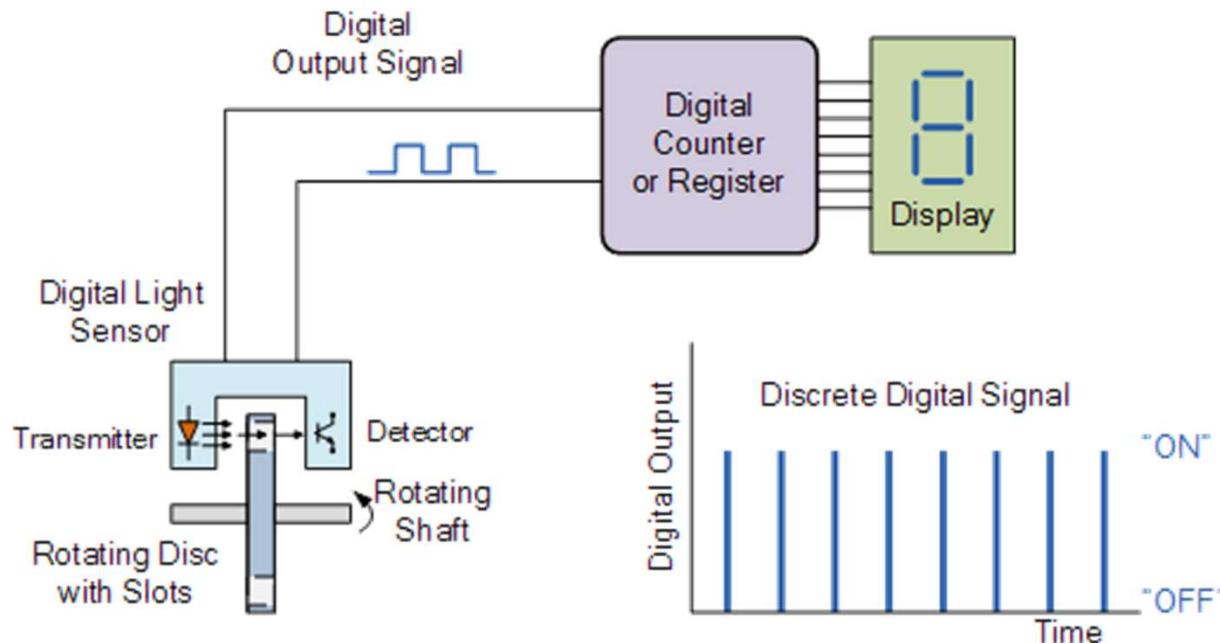
- Analog: Thermocouple used to produce an Analog Signal



[Source: electronics-tutorial.com]

# Digital

- Digital: Light Sensor used to produce a Digital Signal



[Source: electronics-tutorial.com]

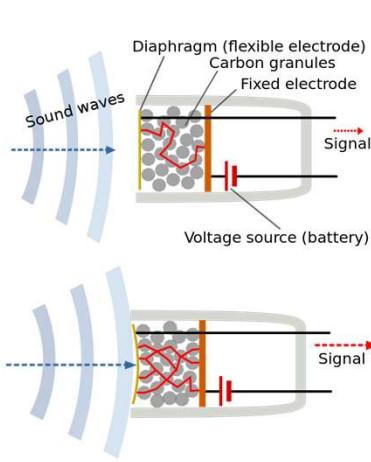
# Sensing what?

---

- Basically everything (see list on Wikipedia...)
  - Acoustic, sound, vibration
  - Chemical (includes all biomedical ones)
  - Electric current, electric potential, magnetic, radio
  - Environmental
  - Flow, fluid velocity
  - Ionizing radiation, subatomic particles
  - Navigation
  - Position, angle, displacement, distance, speed, acceleration
  - Optical, light, imaging, photon
  - Pressure, Force, density, level
  - Thermal, heat, temperature
  - Proximity, presence

# Some examples

- A simple one: a microphone
  - Can be very technological (e.g., a MEMS microphone)



## Features

- Single supply voltage operation
- Fully differential output
- Omnidirectional sensitivity
- High signal-to-noise ratio
- High bandwidth
- Package compliant with reflow soldering
- High RF immunity
- ECOPACK®, RoHS, and "Green" compliant

## Description

The MP23AB01DH is a compact, low-power microphone built with a capacitive sensing element and an IC interface.

The sensing element, capable of detecting acoustic waves, is manufactured using a specialized silicon micromachining process to produce audio sensors.

The MP23AB01DH has sensitivity of 38 dB  $\pm 1$  dB, an acoustic overload point of 135 dB SPL with minimum 65 dB signal-to-noise ratio.

The MP23AB01DH has fully differential output in order to minimize common mode noise.

The MP23AB01DH is available in a package compliant with reflow soldering and is guaranteed to operate over an extended temperature range from -40 °C to +85 °C.

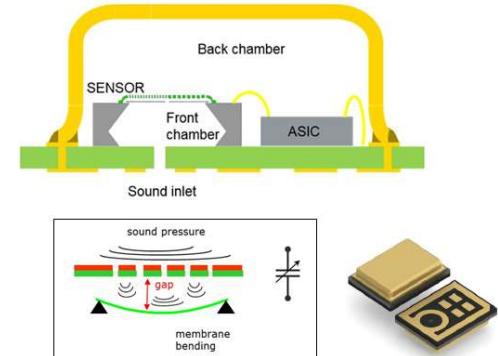
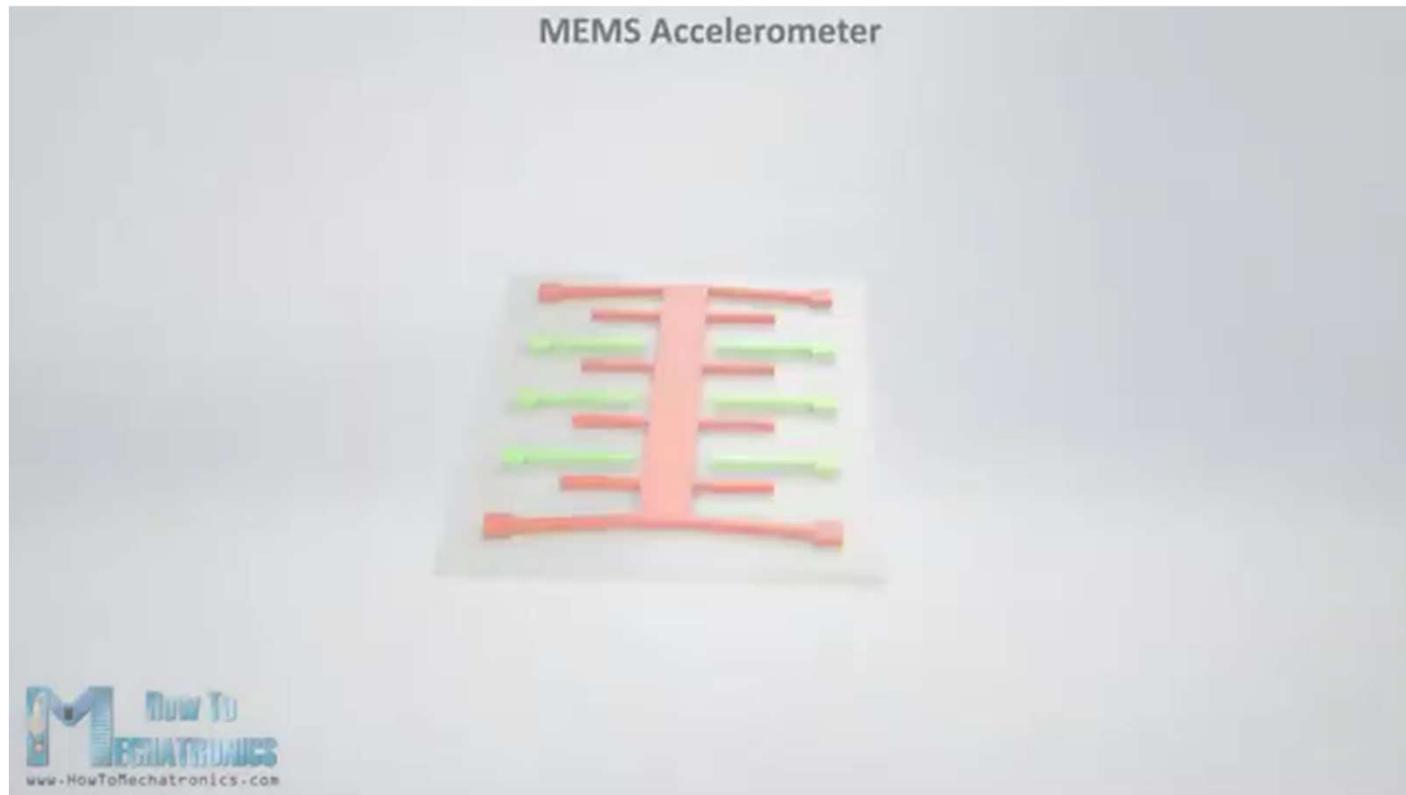


Table 1: Device summary

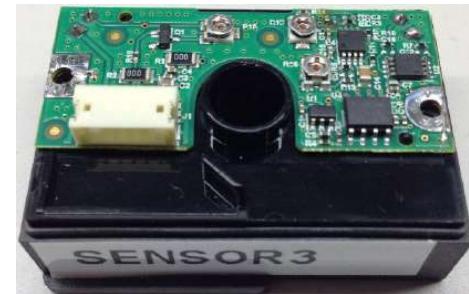
Order code	Temp. range (°C)	Package	Packing
MP23AB01DH	-40 to +85	(3.35 x 2.5 x 0.98) mm	Tray
MP23AB01DHTR	-40 to +85	(3.35 x 2.5 x 0.98) mm	Tape and reel

# MEMS: MicroElectricalMechanicalSystems



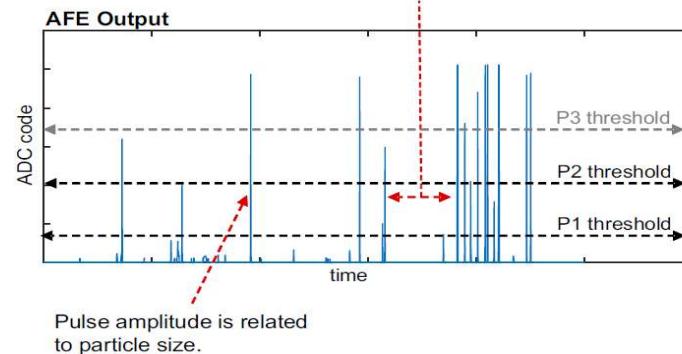
# Some examples

- A more sophisticated one: PM 2.5-10 sensor
  - E.g., for smoke, dust, pollen
  - Works by detecting with a photodiode the light scattered by particles suspended in air in a chamber
  - A circuit generates a series of pulses that are directly related to the light scattered by the particles
  - Needs a lot of processing!

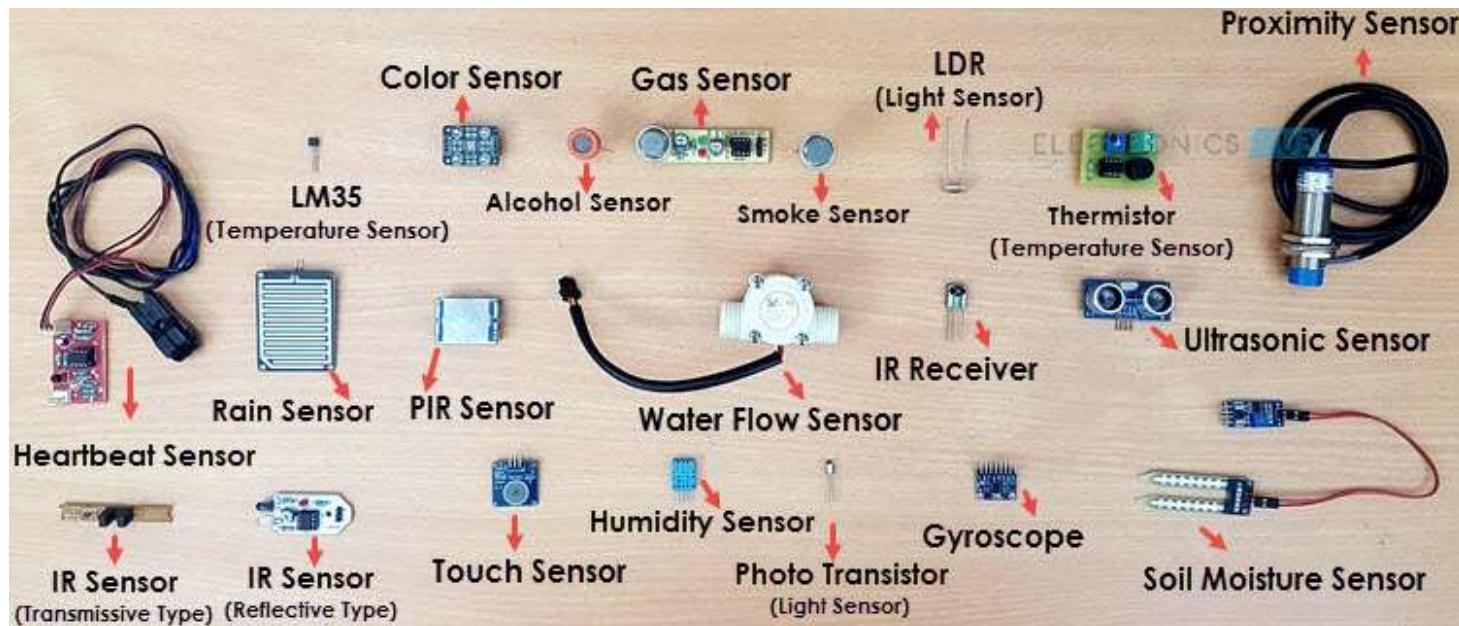


TIDA-00378 PM2.5 and PM10

Pulse spacing (frequency) is related to particle concentration



# And many others...



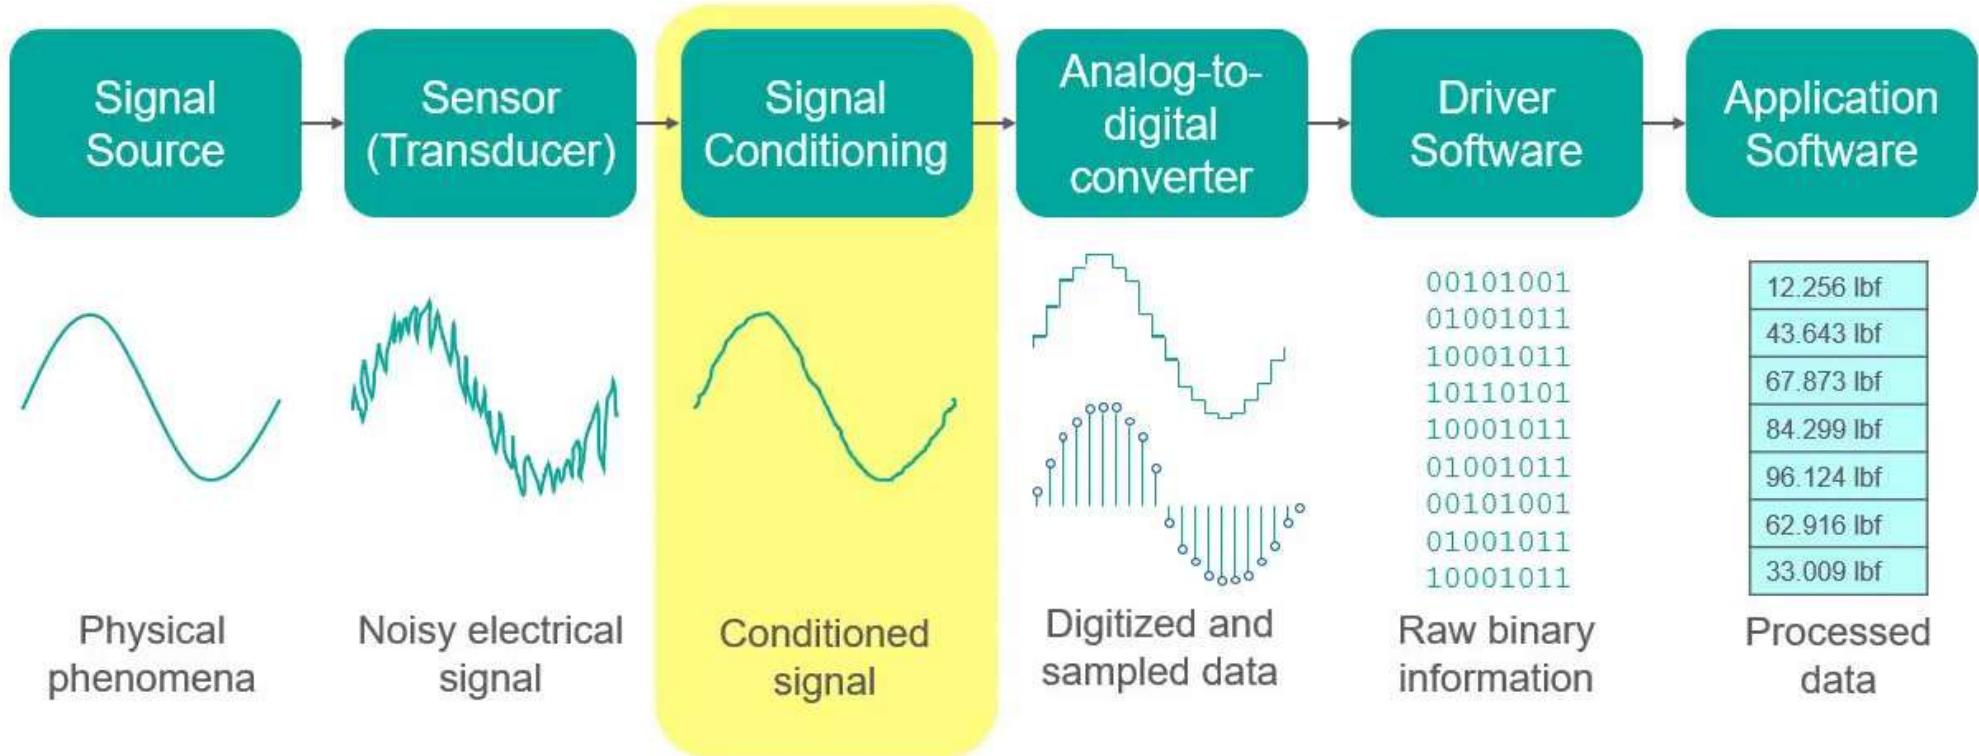
# Concept generalization

- Key element is the transducer
- Transducer = element that converts energy from one domain to another



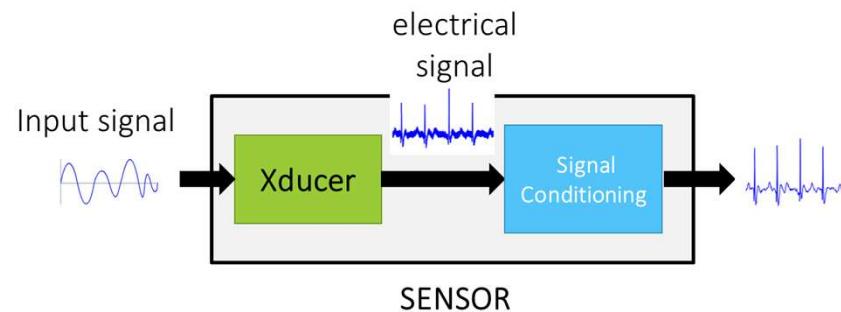
- Typically we are interested into transducers that have electrical domain as output
- A sensor (as well as an actuator) includes a transducer
  - Sometimes the same actuator can serve in both directions (e.g., an antenna)

# From the physical world to a digital domain



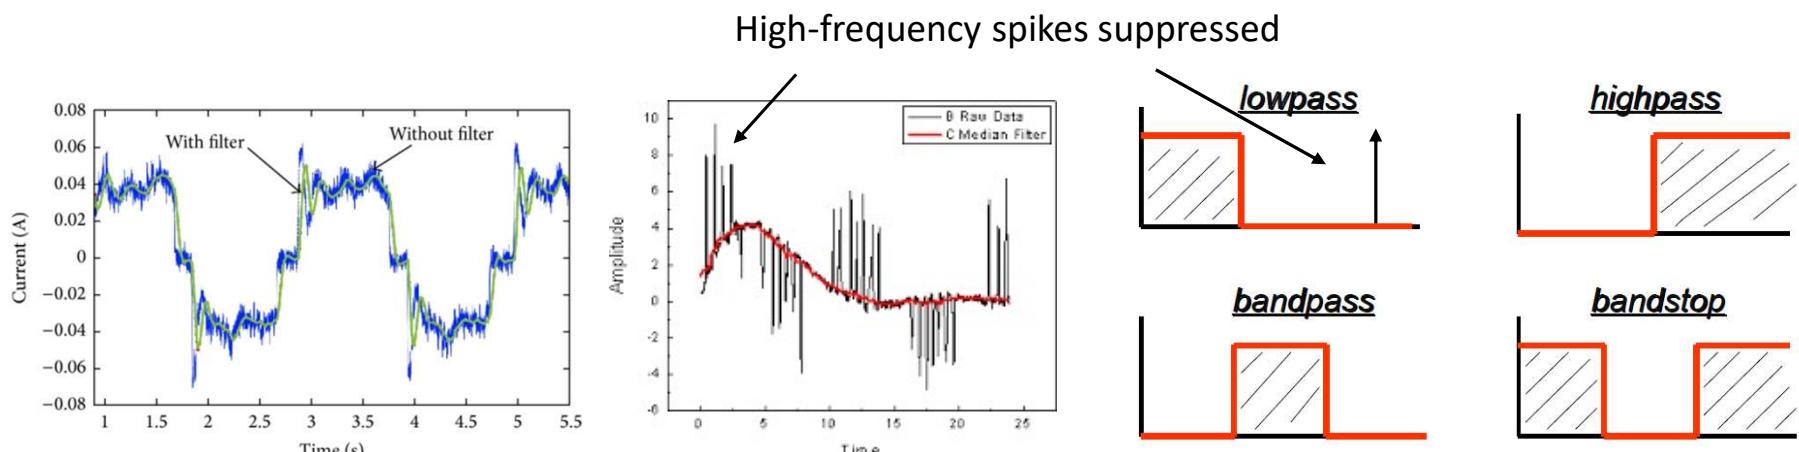
# Signal conditioning

- Adjust the dynamic of signals generated by the transducer in order to match the desired electrical levels
- Still in the analog domain
  - Analog circuits
- Part of any off-the-shelf sensor
  - Active/Passive
  - Digital/analog



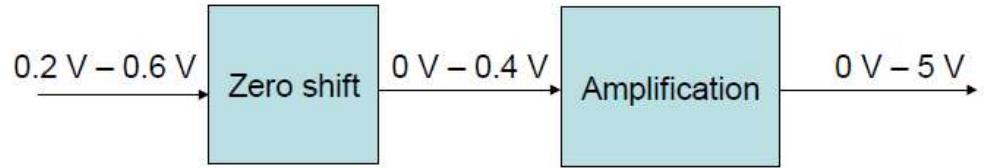
# Signal conditioning

- Filtering:
  - the most common signal conditioning function, as usually not all the signal frequency spectrum contains valid data.
  - For noise suppression and/or frequency selection (need some specific band of the signal only)

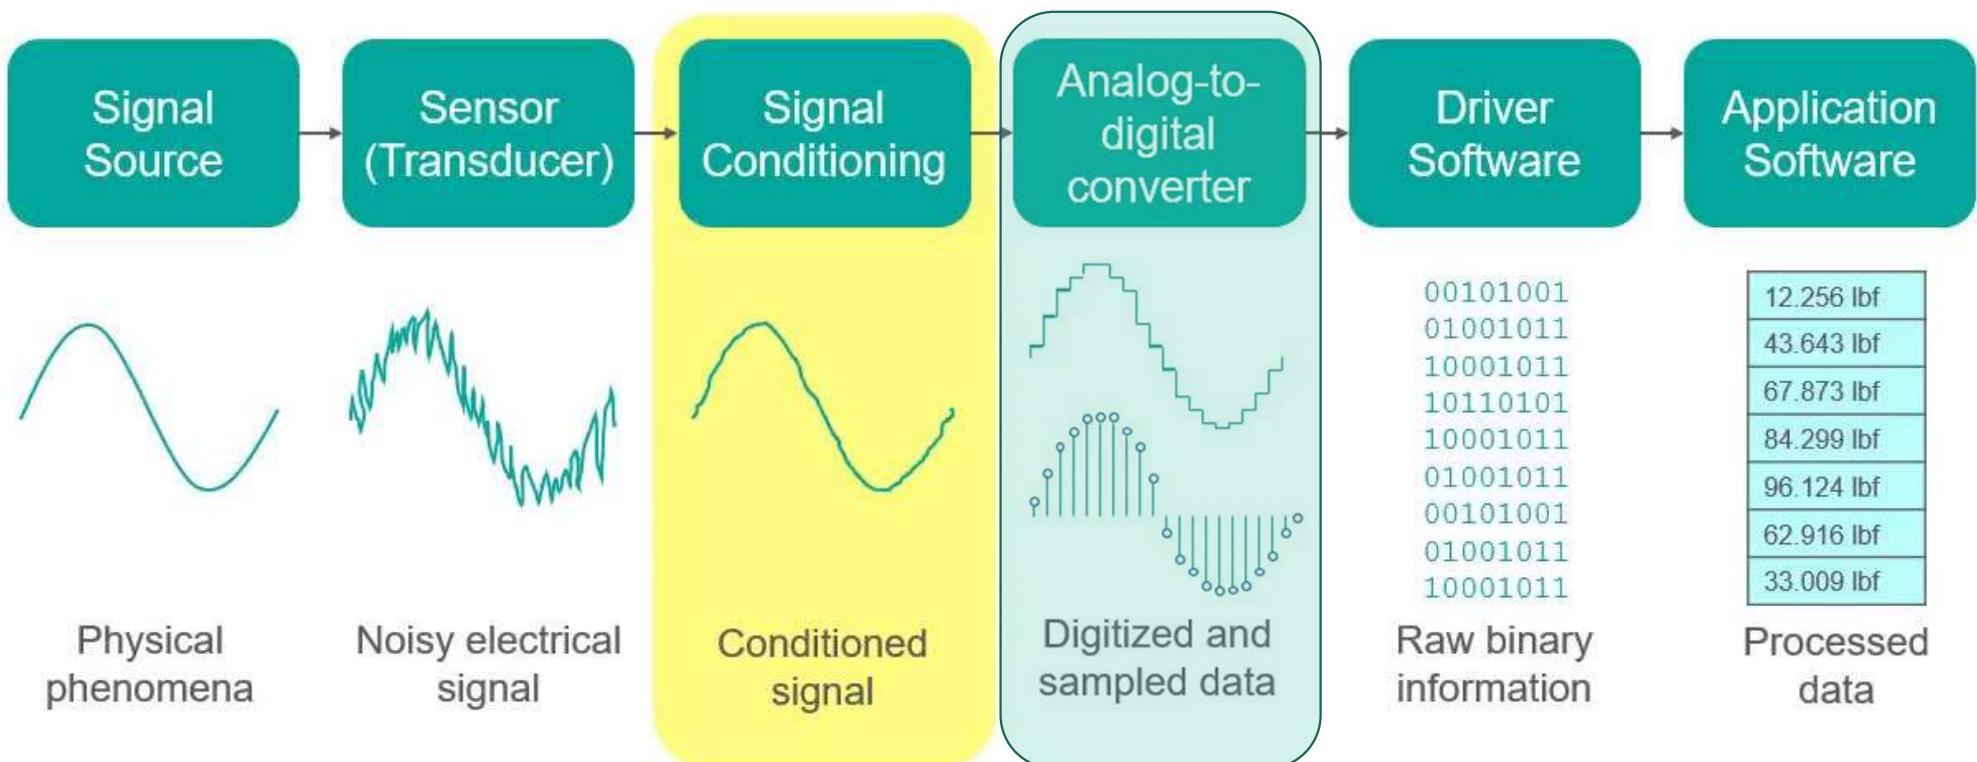


# Signal conditioning

- Level shift
  - adjust the bias (zero value)
  - Better exploit the dynamic of ADCs
- Amplification:
  - Signal amplification performs two important functions:
    - increases the resolution of the input signal
    - and increases its signal-to-noise ratio
  - Signal levels are usually in the  $\mu$ Vs to mVs range!
    - For example, the output of an electronic temperature sensor, which is probably in the millivolts range is probably too low for an analog-to-digital converter (ADC) to process directly. In this case it is necessary to bring the voltage level up to that required by the ADC. [wiki]

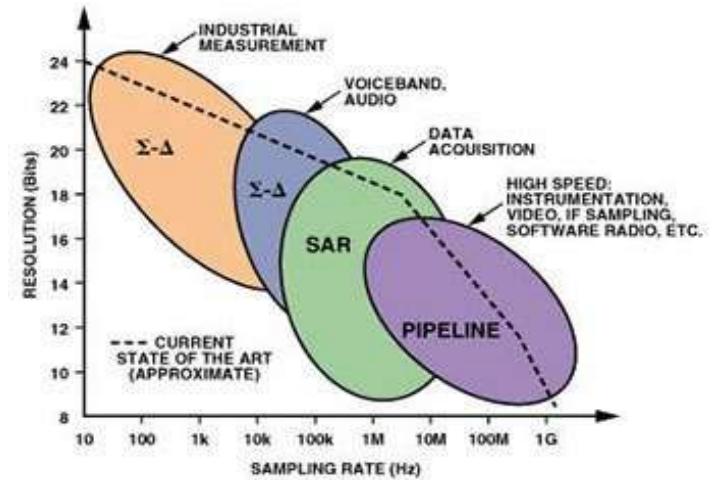


# Converting to digital



# Converting to digital

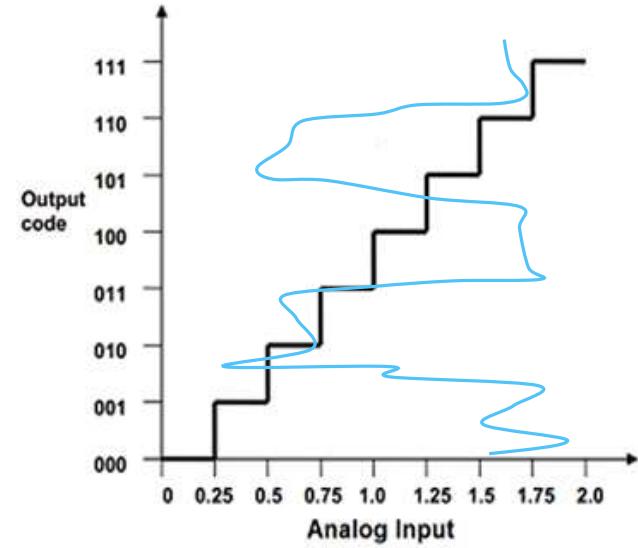
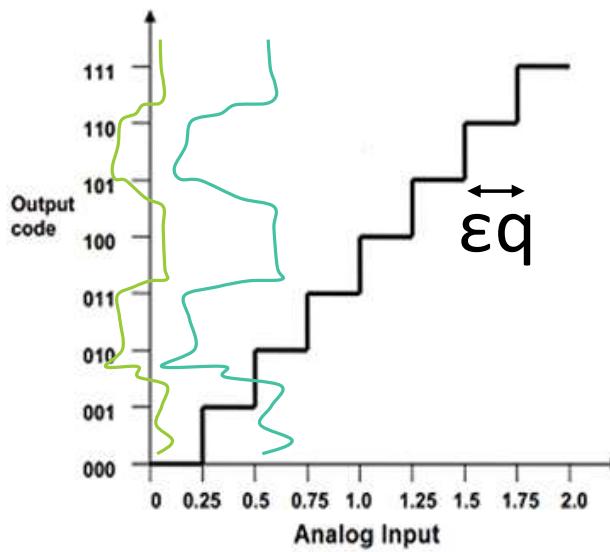
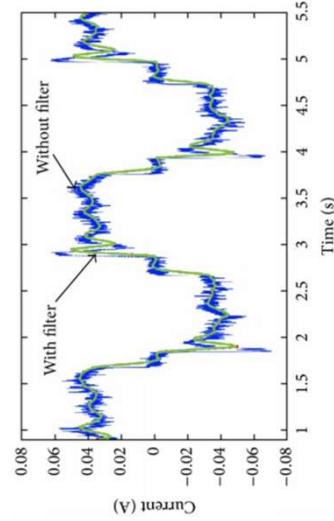
- Regardless of the sensor type the conversion into the digital domain is mandatory
  - Processing IS digital!!!
- A/D conversion implements sampling + quantization
- Key parameter is the resolution, i.e., # of bits of the quantization
- Various circuit implementations!



Source: Analog Devices

# ADC operations

- Conceptually, an ADC implements a step-wise transfer function
  - Example: 3-bit conversion
- cover the whole input dynamic = minimize the quantization error
  - Conditioning is key!



# ADC Example

---

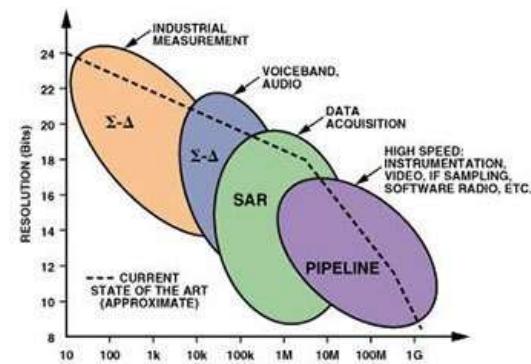
- 10-bit ADC
  - Input voltage  $x$  [0-5V]
  - 0V is 0, 5V is 1023, i.e.  $2^{10}-1$
  - $1023:5 = Y_{bin} : x$
- An analog voltage  $x$  [Volts] will be reported by the ADC as

$$Y_{bin} = \text{ceil}(x * 1023/5)$$

- Example: input  $V=2.12V$ 
  - $Y_{bin}=\text{ceil}(433.752) \rightarrow$  reading is 434
  - $\epsilon = 0.00121 V$  (0.06%)

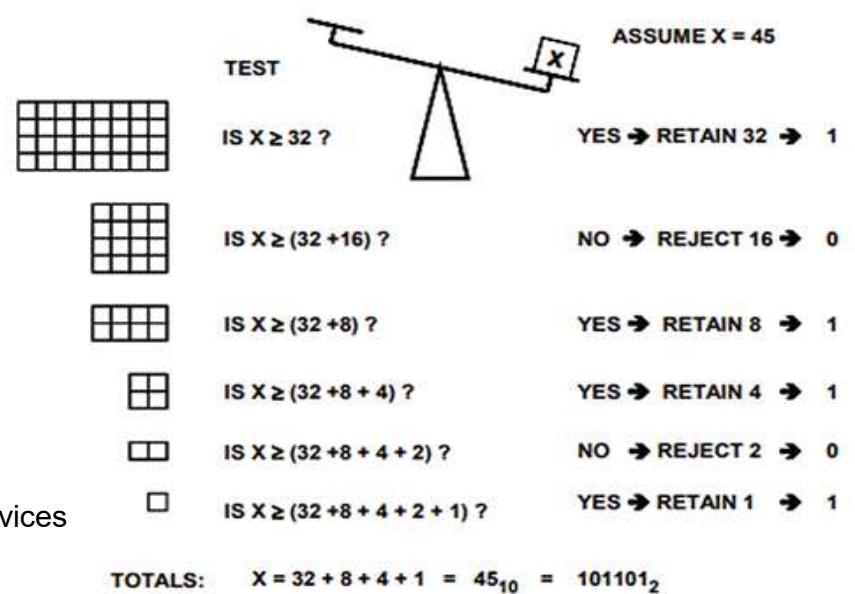
# ADC operations

- Many implementations do exist
- Successive Approximation Registers (SAR) ADCs
  - Basic principle of operations (iterative!):

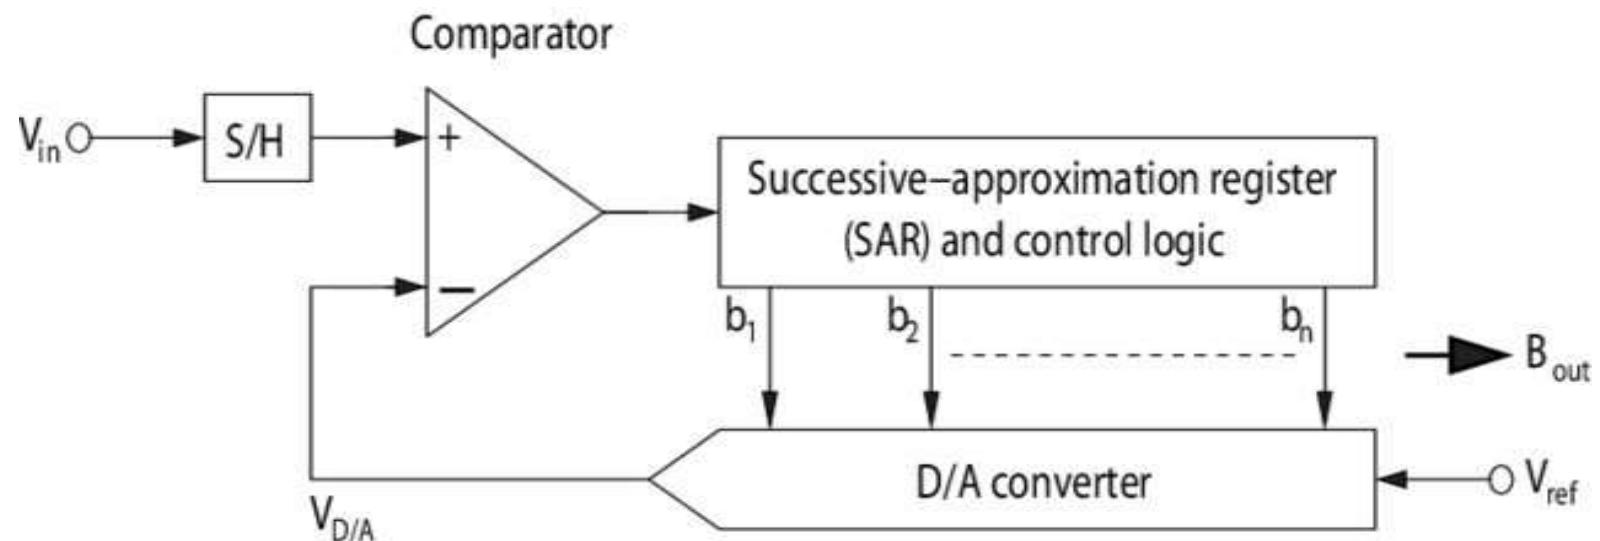


Source: Analog Devices

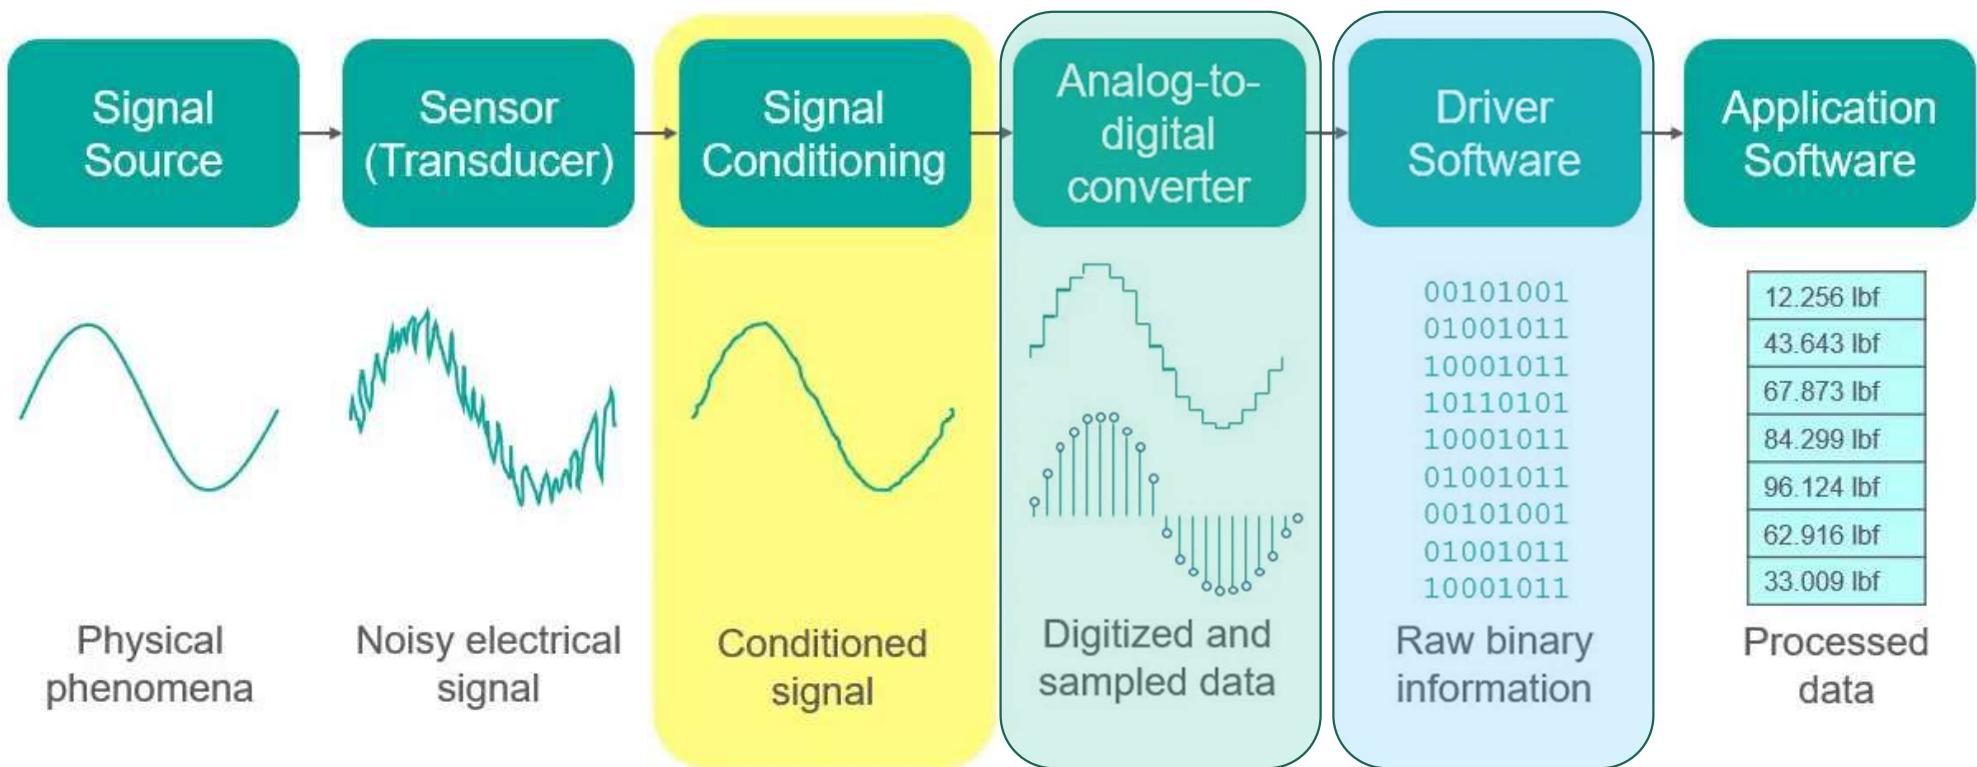
Source: Analog Devices



# (very) Basic SAR architecture



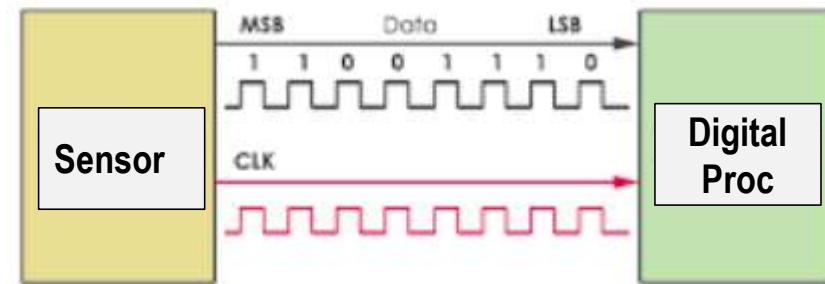
# Readout



# Reading

- Sensors (i.e., ADCs) typically outputs bits **serially** (one at a time):

- Fewer pins
- Better signal robustness (no interference)
- Faster speed at iso-noise



- In the digital world most interfaces are serial!
  - Most popular: USB, SPI, I2C
- This process is transparent to the programmer
  - Drivers manage the transfer protocol
  - A simple read function provides the full word

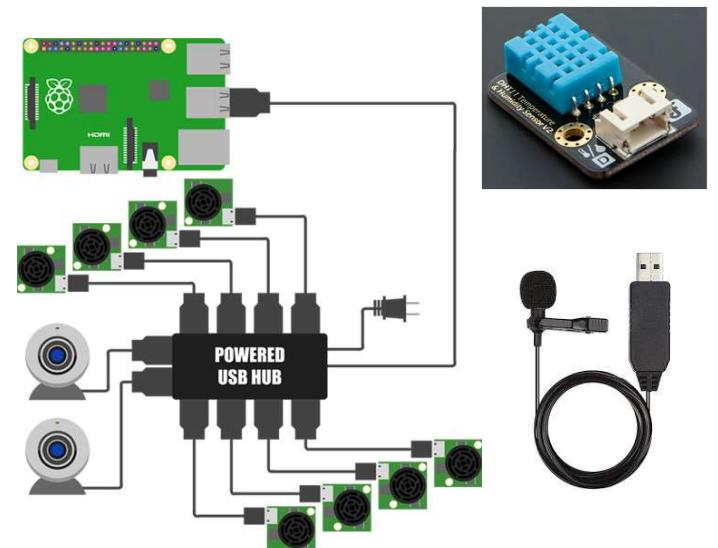
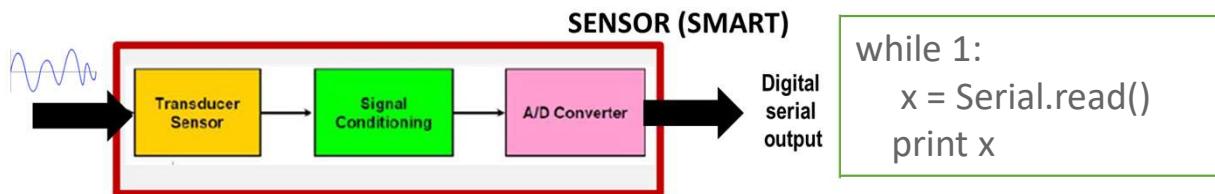
# Serial Interface

---

- USB
  - Vast majority of modern devices (less popular for sensors)
  - A (host/master) B (peripheral/slave)
    - OTG: "On The Go": Accepts A (make me a host) or B (make me a peripheral)
  - Often used as power source
- I2C (Inter-Integrated Circuit)
  - Allows several "slave" (sensor) devices to be accessed by same "master"
  - Data (SDA) and clock (SCL) lines (2 wires)
- SPI (Serial Peripheral Interface)
  - Multiple slaves, bidirectional
  - 4 wires (Clock, select, MISO/MOSI)

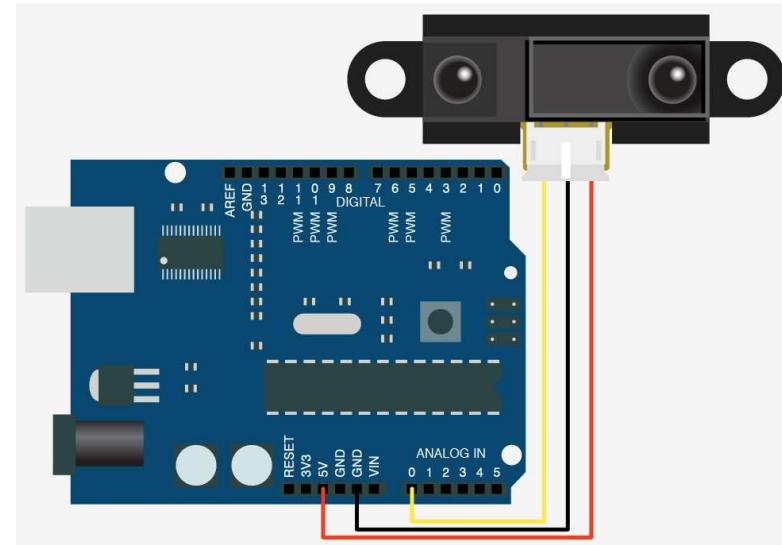
# Practical considerations

- How can a non-EE engineer truly use sensors?
  - Too much technicalities, interfaces, etc.
- Commercial (smart) sensors for system designers include all the components and hide all the details to the designer
  - also preventing possible optimizations...
- Smartness often include
  - Full analog-to-digital pipeline
  - Programmability (signals to set some parameters)
  - Storage (e.g. for config parameters or buffering)



# Practical considerations

- Simpler sensors may demand the ADC conversion to “someone else”
  - done on-board and not in-sensor



# Coordinating multiple sensors

---

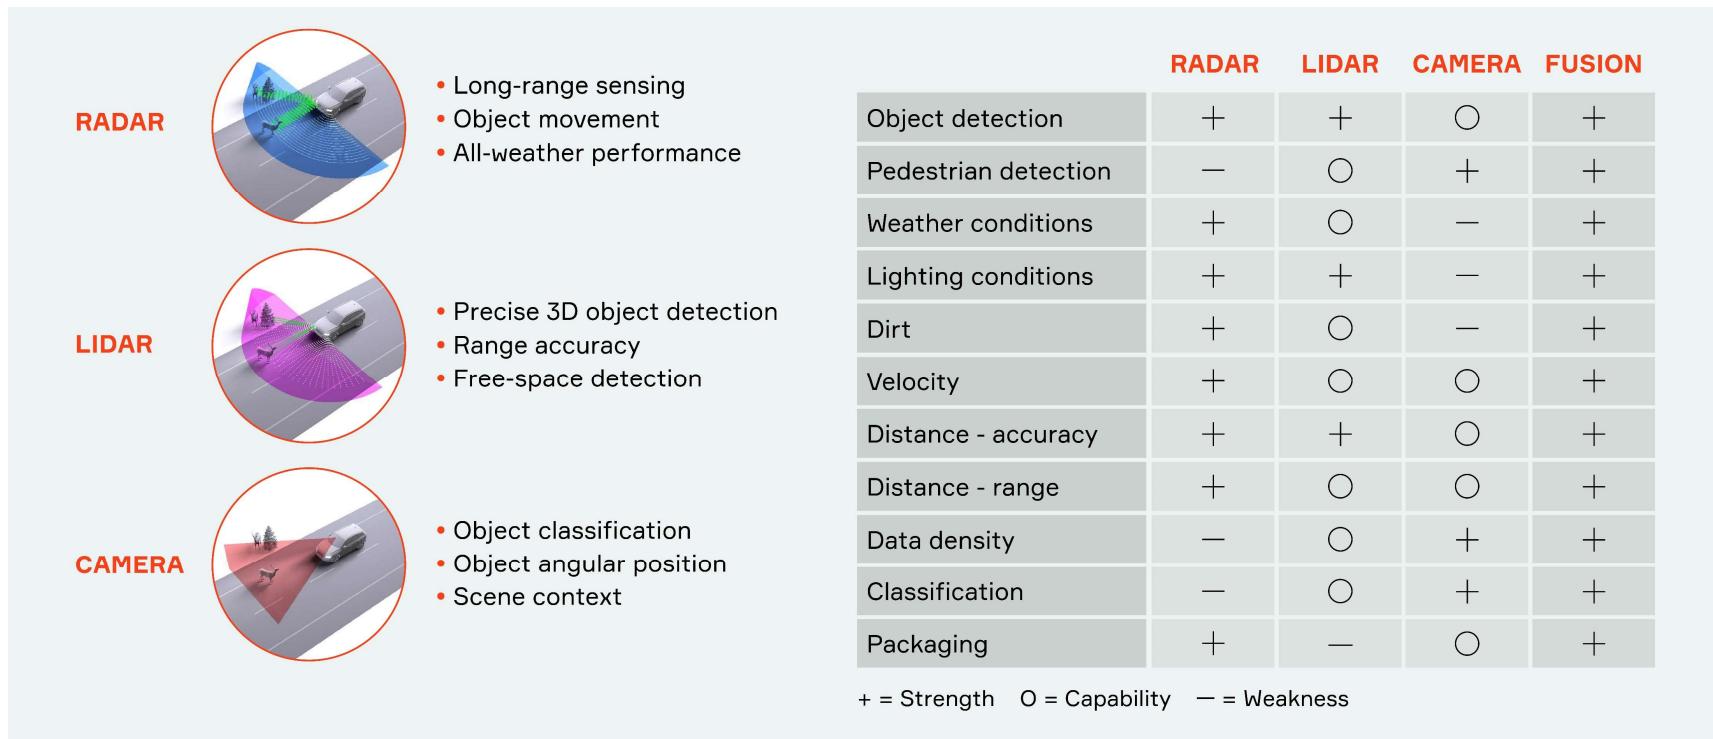
- Devices seldom have only one type of sensor
- Multiple sensors can be used to sense independent quantities for
  - uncorrelated usage/multiple functions
    - E.g., a camera and microphone serve different purposes...
  - correlated usage
    - To improve the quality of the sensing process
      - E.g., gyroscope and accelerometer to improve an estimate of motion
    - Called sensor fusion
- Sensor fusion: combining sensory data so that the resulting information has less uncertainty than would be possible when these sources were used individually
  - Information integration problem

# Why sensor fusion?

---

- Accuracy & reliability
- Fused data from multiple sensors provides several advantages over data from a single sensor:
  - Statistical advantage:
    - adding  $N$  independent observations the estimate of the target measure is improved by a factor proportional to  $N^{1/2}$ , assuming the data are combined in an optimal manner
    - NOTE: The same result could also be obtained by combining N observations from an individual sensor
  - Better observability:
    - Unlike the case of a single sensor, where some quantities cannot be accurately estimated

# Self-driving

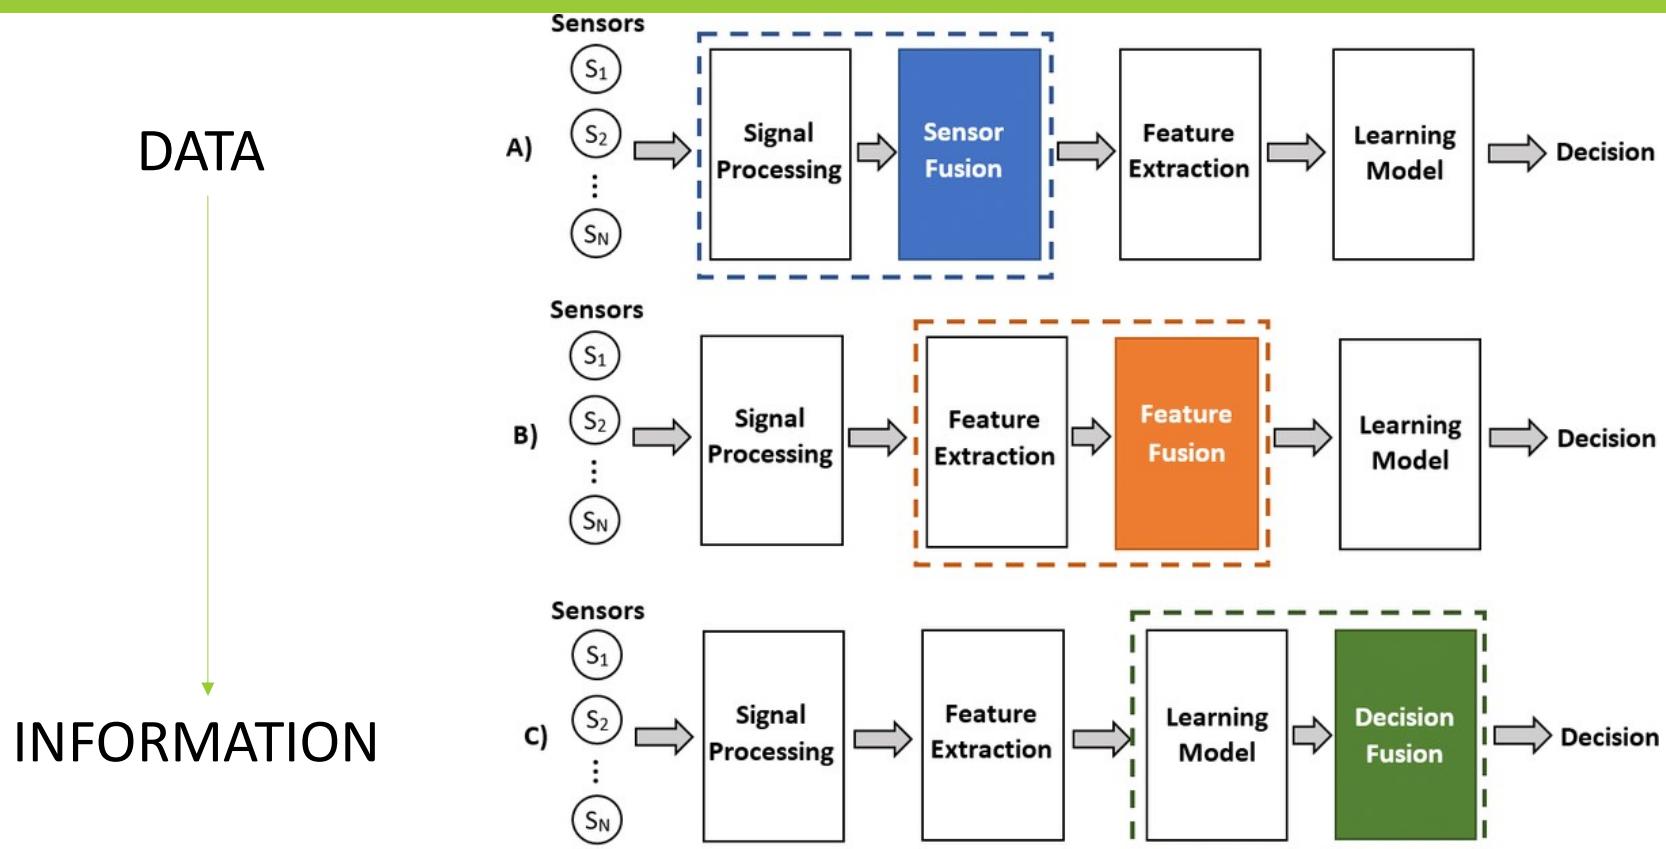


# Sensor Fusion

---

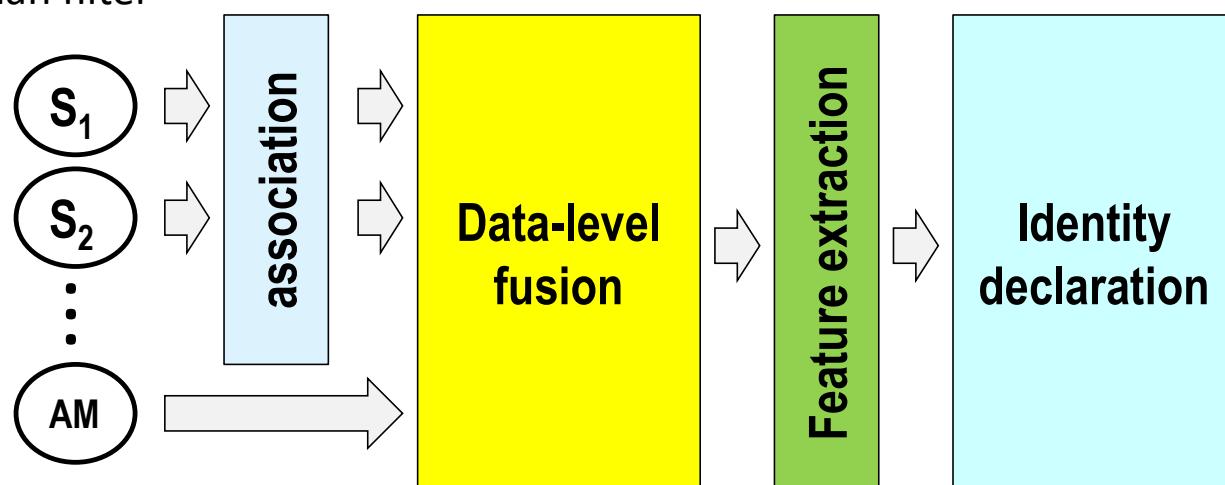
- There are three fundamental paradigms
  - A. Measurement fusion: Direct fusion of sensor data
  - B. Feature fusion: Representation of sensor data via feature vectors and fusion of the feature vectors
  - C. Decision fusion: Processing of each sensor to achieve high-level inferences and fusion of these inferences

# Sensor Fusion



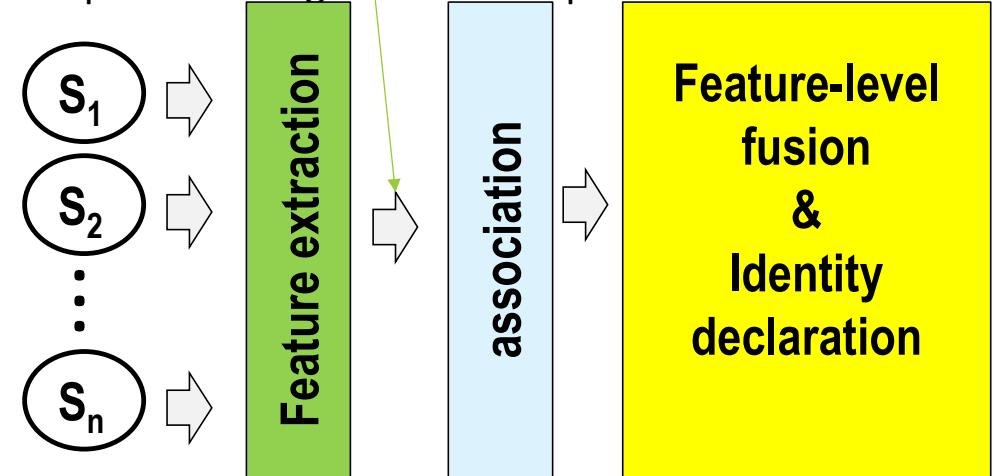
# A. Measurement fusion

- Suitable when data are commensurate
  - sensors are observing the same physical phenomena, eventually measuring different physical quantities
- Classical estimation models (control theory)
  - Kalman filter



## B. Feature fusion

- For non-commensurate sensor data, data **must** be fused at the **feature** level
  - Features are combined into a single concatenated feature vector
- Feature: a sort of signature/digest representative of the sensed quantity
  - Works as a common ground
  - features are extracted from multiple sensor observations and combined into a single concatenated feature vector that is input to pattern recognition techniques such as neural networks, clustering algorithms



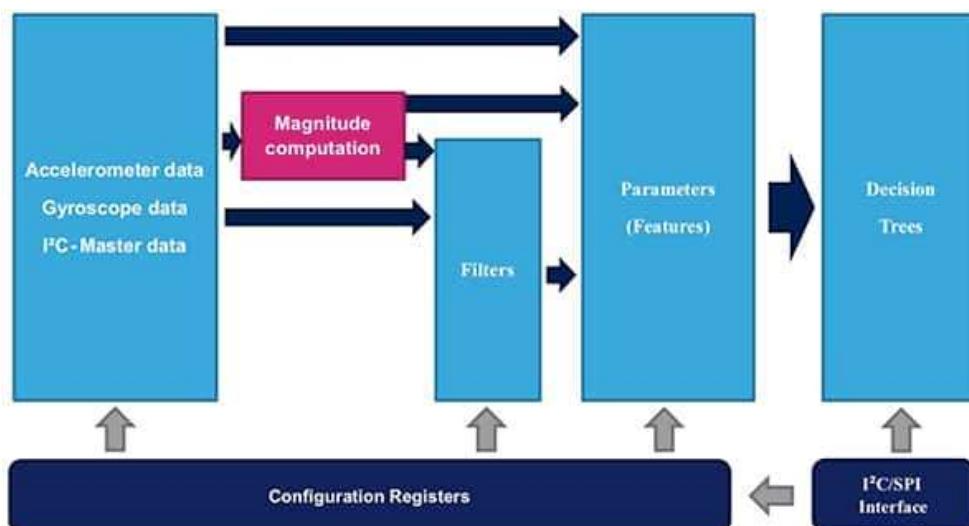
# More about features

---

- Features are strongly sensor- and application-dependent!
- Since they must be combined, they are often normalized quantities
- Some examples:
  - Sound features
    - Pitch, Centroid, Spectral variance, Zero-crossing rate (ZCR). Root Mean Square (RMS) (Volume), Signal-to-noise ratio (SNR), ...
    - Suitable also for generic signals
  - Visual (image) features
    - Edges, corners, ridges, color, hue, blobs, Spectral properties, special filters, etc

# A practical example: Smart Motion from STM

- Motion sensor

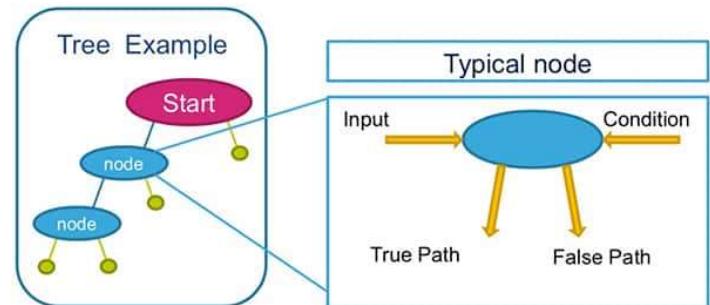


- the device generates a large number of features calculated from primary data within a sliding time window including:
  - mean
  - variance
  - energy
  - peak to peak
  - zero crossing
  - positive zero crossing
  - negative zero crossing
  - peak detector
  - positive peak detector
  - negative peak detector
  - minimum
  - maximum

# A practical example: Smart Motion from STM

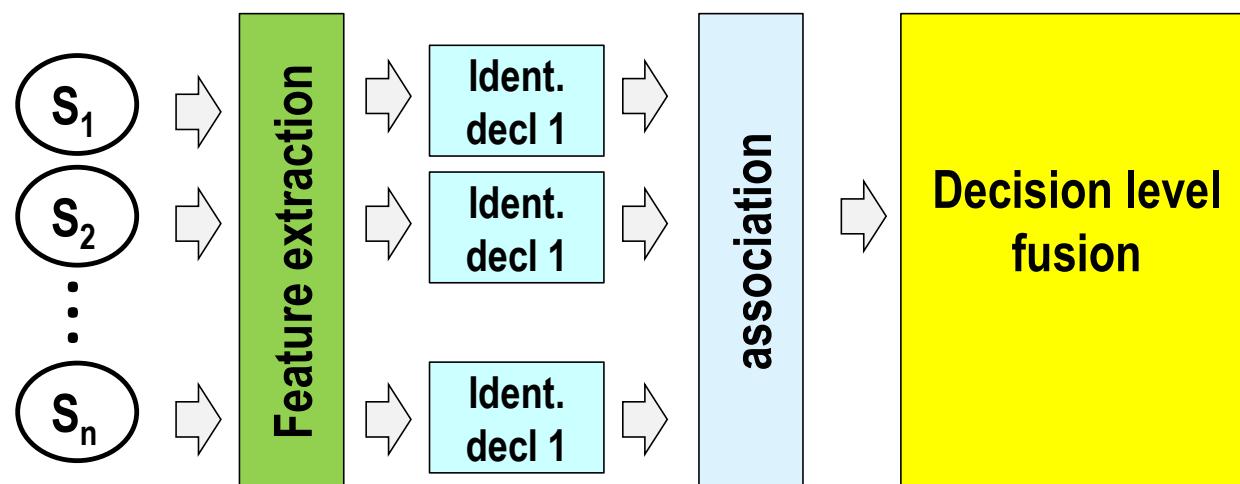
- at each update cycle, the decision tree would be invoked to work through its nodes to determine if the data available – with that update – represented no movement, forward movement, or some other movement as follows:
  1. test the magnitude of an accelerometer measurement
    - 1.1. terminate if the value is below some predetermined value (the condition)
    - 1.2. otherwise, branch to a child node to test gyroscope measurements taken in the same time window
      - 1.2.1. terminate if gyroscope measurements are below some predetermined value or
      - 1.2.2. continue to a deeper child node to test other attributes measured in the same time window or test the same attribute against another condition.
  - This process repeats until the test reaches a terminal node, which in this context corresponds to a particular complex motion event, or class. In this simple example:
    - terminal node 1.1 might indicate that the data, or feature set, should be classified as “no movement”
    - terminal node 1.2.1 might indicate that the feature set should be classified as “forward movement”
    - terminal nodes below node 1.2.2 might be indicative of a moving turn or more complex change in orientation

Feature fusion using decision tree

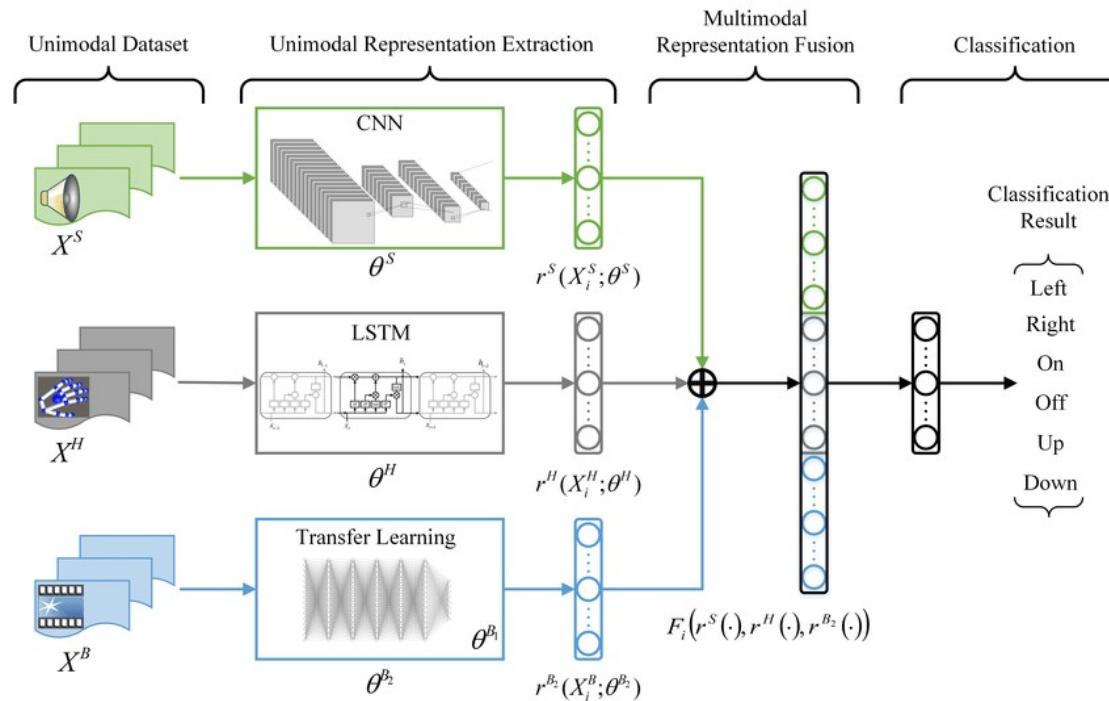


## C. Decision fusion

- Each sensor has made a preliminary determination of an entity's location, attributes and identity before combining
- Decision-level fusion algorithms are used, such as
  - Voting/weighted decision or more complex algo (Bayesian inference and Dempster-Shafer's method)



# A practical example



# Other considerations

---

- Sensor fusion is a computational issue
  - Consider it as a part of the “software” part of an IoT device
  - Future smart sensors will integrate it as an embedded feature
- It requires significant computation complexity
  - Feature extraction and in particular classification are complex algorithms!
  - Tradeoff between quality of final classification and complexity of the method
- In-sensor computing will soon become a cutting-edge technology

## Actuators

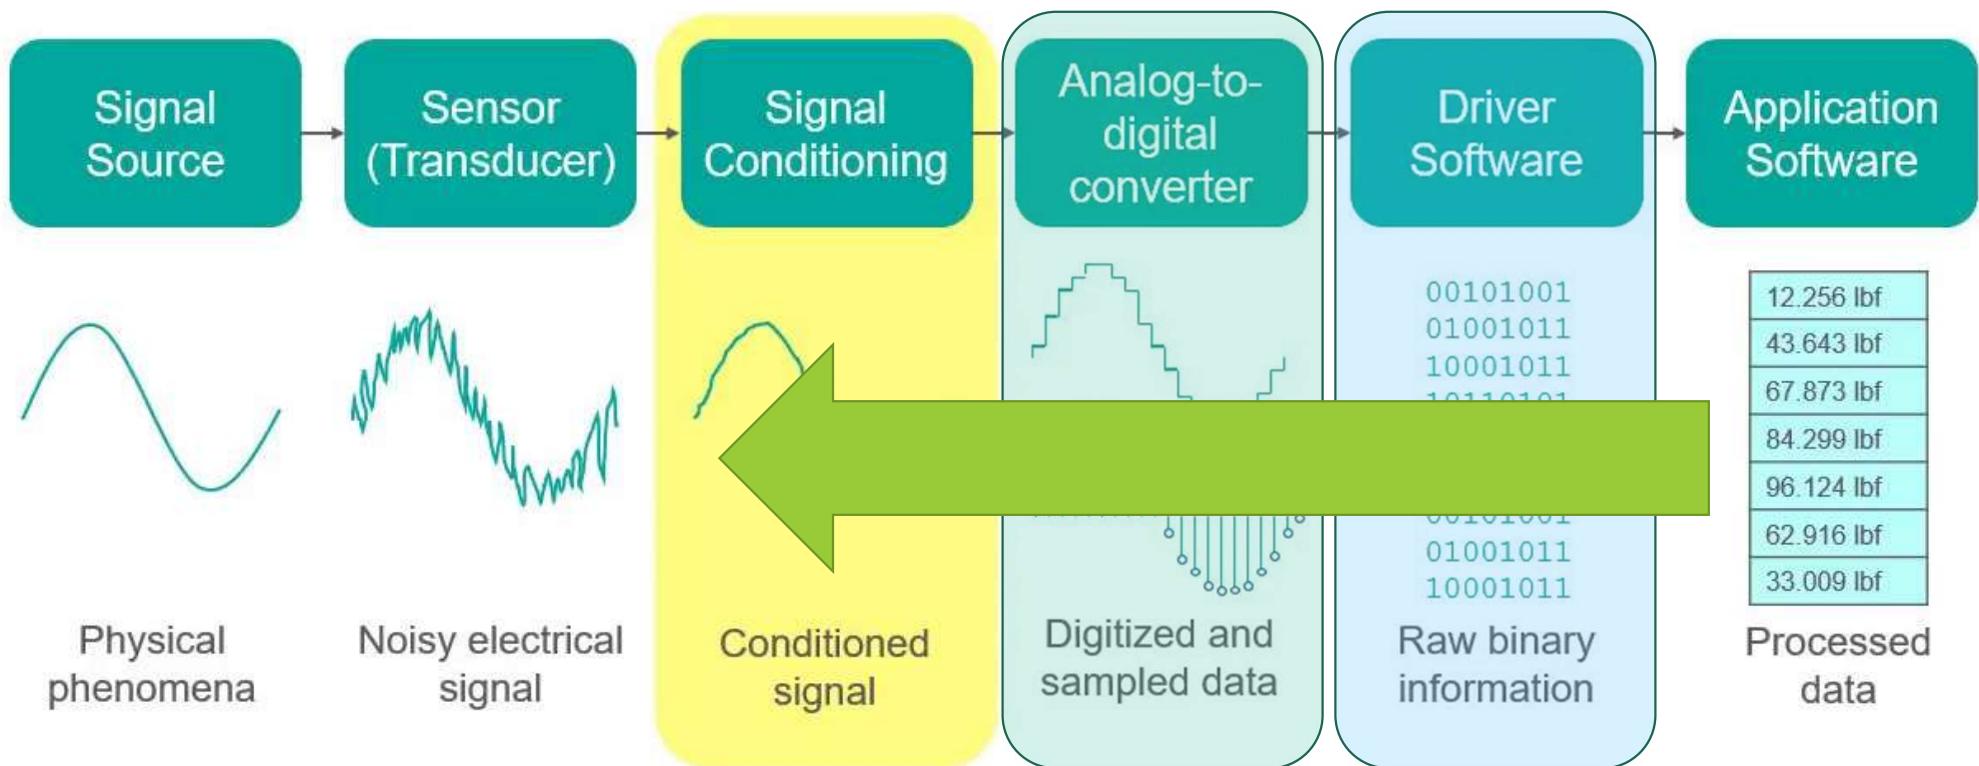
---

# Actuators: back to the analog world

---

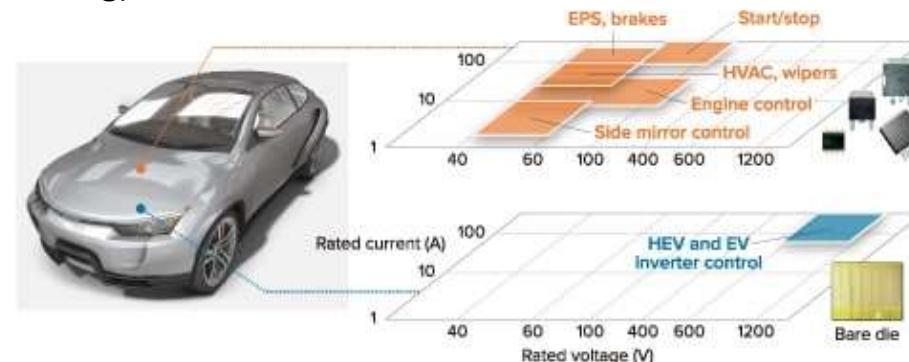
- Actuator = devices that alter a physical quantity
- The symmetric action of sensing
  - DAC + transducer
  - Transducer : electrical domain is in input...
- As for sensors, a number of possible devices
  - Electric motors
  - Comb drives
  - Piezoelectric actuator
  - Servomechanisms
  - Pneumatic actuators
  - Hydraulic actuators
  - Solenoids
  - Stepper motor
  - Speakers?

# Backward flow



# Actuators need power!

- Mechanical parts (actuators in general) need more power to be “driven”
  - Digital cores work in the range of hundreds of mV and uA (at best)
- Example: driving a speaker
  - Speaker essentially a resistor (typical  $8\Omega$ )
  - Power rating e.g. 1W  $\Rightarrow$  current  $P=RI^2 \Rightarrow I = \sqrt{P/R} = 350\text{mA}!!!$
  - Assuming MCU has analog output, output current in the range of tens of mA!
    - Need an amplifier (sensor conditioning)...
- Think about electrical motors
  - Power rating KW
  - Need high-power electronics



**Power MOSFETs**

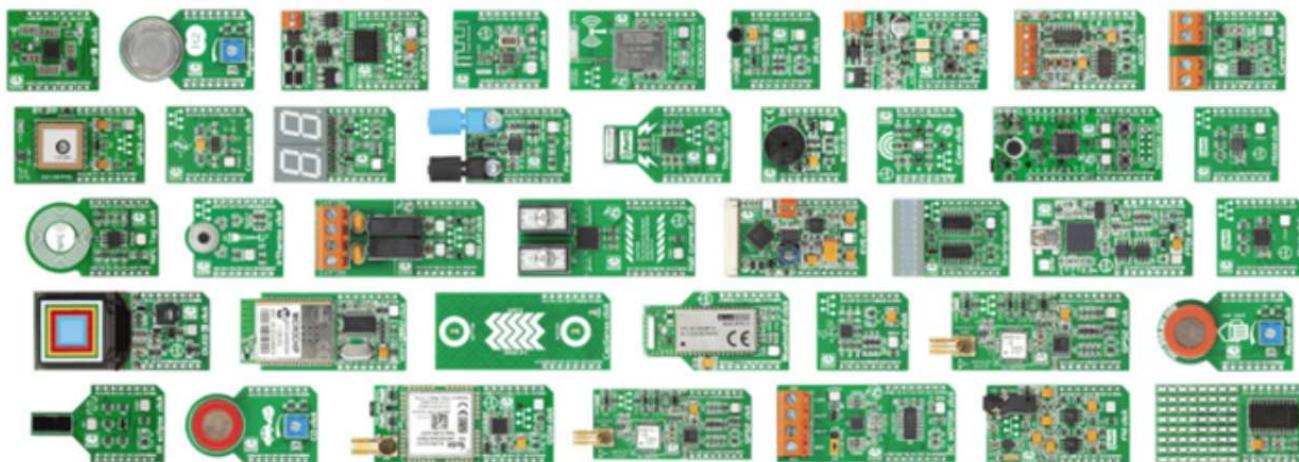
- World-top-class performance
- Established production track record and high quality
- Support for new applications/48V applications

## IGBT

- Development of high-performance IGBTs for HEV/EV propulsion inverter applications
- Voltage rating variations from 650V to 1,200V

# Custom interfaces

- Apart from very common cases, actuators need custom circuitry for the interface
- Some integrated solutions do exist...
  - Rely on a common, standardized system architecture that allows seamless plug&play of sensors
  - Example: **MikroBUS™ Adapter click boards™**

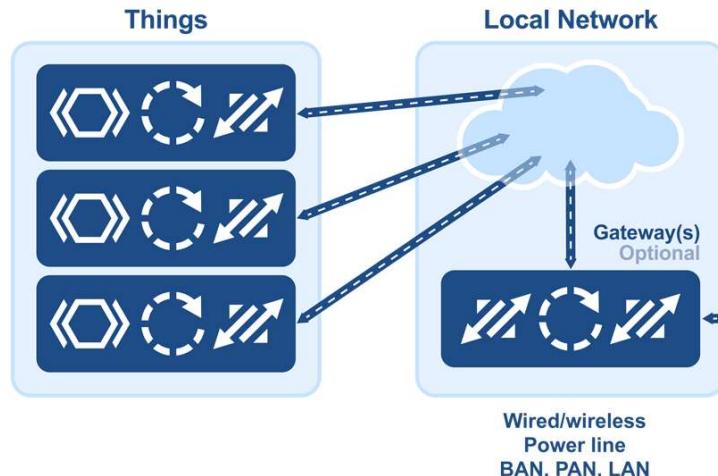


Processing

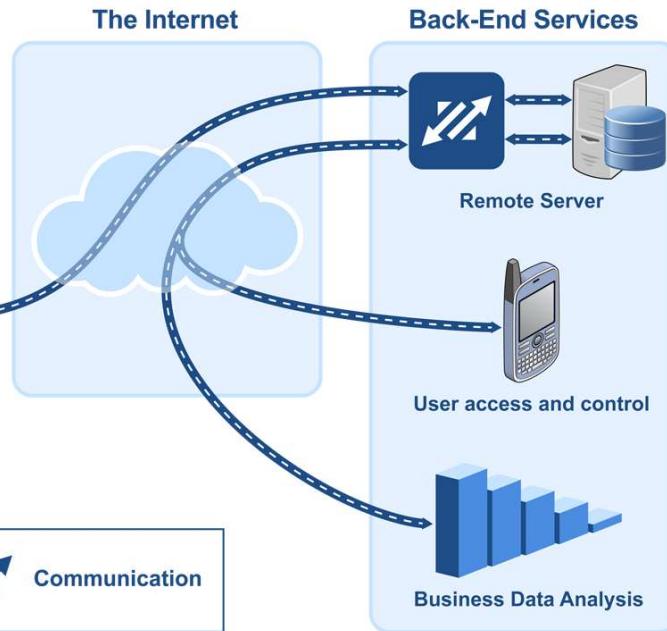
---

# Different kind of “processing”

MCU: (Micro-)controllers



MPU: (micro-)processors



# MCU vs MPU

- MPUs contain only a CPU (possibly with some cache) and interfaces
  - They require added peripherals to perform tasks.
- MPUs are quite **powerful**
  - Computational power (multiple execution units (multicore))
  - Support for special operations
  - Accelerators
- MPUs are ‘**modular**’
  - Resources (memory, disk, etc) are external,
  - Can be added as needed
- **Expensive!!!**



**Intel Core i7-7820X**

8 offerte: € 475,10 – € 641,46 \*

0 opinioni: [Lascia un'opinione!](#)

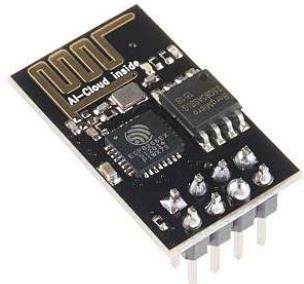
Opinione degli esperti: ★★★★☆

3 Recensioni: **Voto medio 8/10**   ...

[CPU 8 Core](#) · Socket Socket 2066 · Numero di Threads 16 ·  
Frequenza di clock 3,6 GHz · Frequenza turbo massima 4,5 GHz ·

# MCU vs MPU

- MCUs are “complete systems” on a single board or chip
  - They contain RAM, ROM, and similar peripherals
  - A number of interfaces typical of sensors or other devices
  - Often include one or more ADC for analog inputs!
- MCUs have limited functionalities and modest computational power
  - things are changing...
- They are meant for application-specific systems
  - things are changing...
- Cheap!!



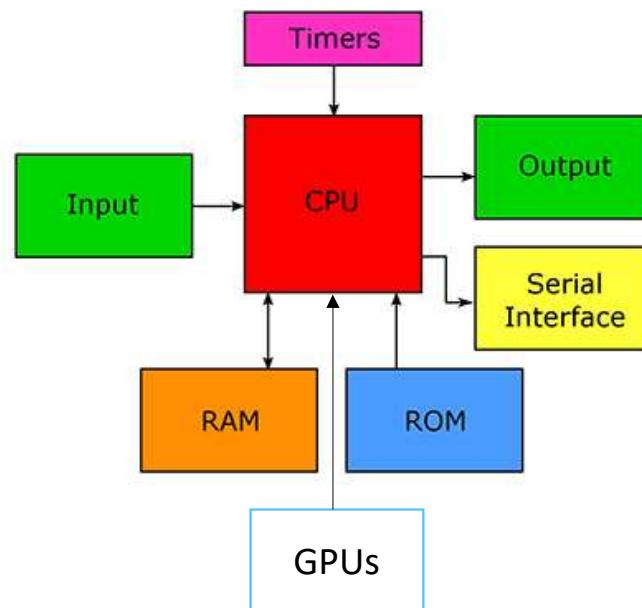
Processor: L106 32-bit RISC running at 80 MHz  
Memory: 32 KiB instruction RAM, 32 KiB instruction, cache RAM, 80 KiB user-data RAM  
External QSPI flash: up to 16 MiB is supported  
IEEE 802.11 b/g/n Wi-Fi Integrated  
16 GPIO pins  
SPI: I<sup>2</sup>C (software implementation), I<sup>2</sup>S interfaces with DMA (sharing pins with GPIO), UART on dedicated pins, 10-bit ADC (successive approximation ADC)

# Some figures

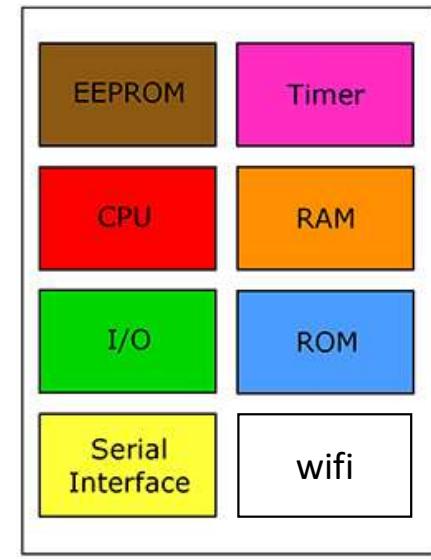
	MPU	MCU
Cost	100's \$	1-10 \$
Memory (RAM)	A few GBs	A few MB
Persistent Storage	Disk, TBs	Can be absent, otherwise FLASH or ROM (kB)
Clock	GHz	KHz/MHz
Cores	2-16	1
Parallelism	32-64	4-32
Speed	100-200 GOPS	1-10 MOPS
Power	100s of W	10-100 uA/MHz (a few mWs @ max f)

# Conceptual Block diagram

Microprocessor: CPU  
and several supporting chips.

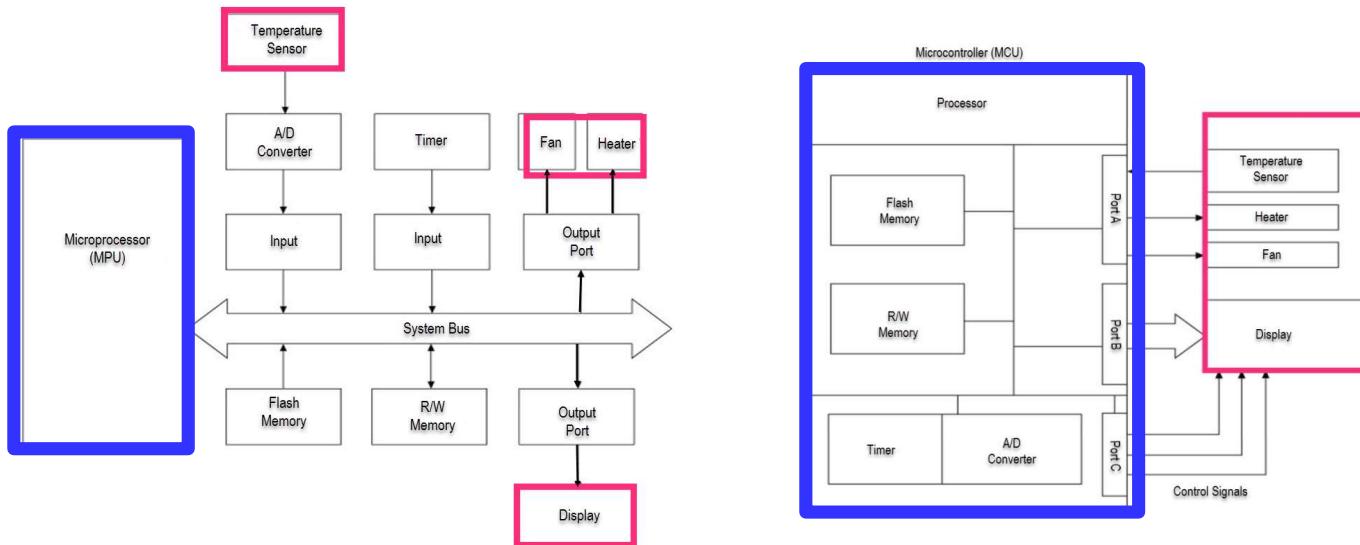


Microcontroller: CPU  
on a single chip.



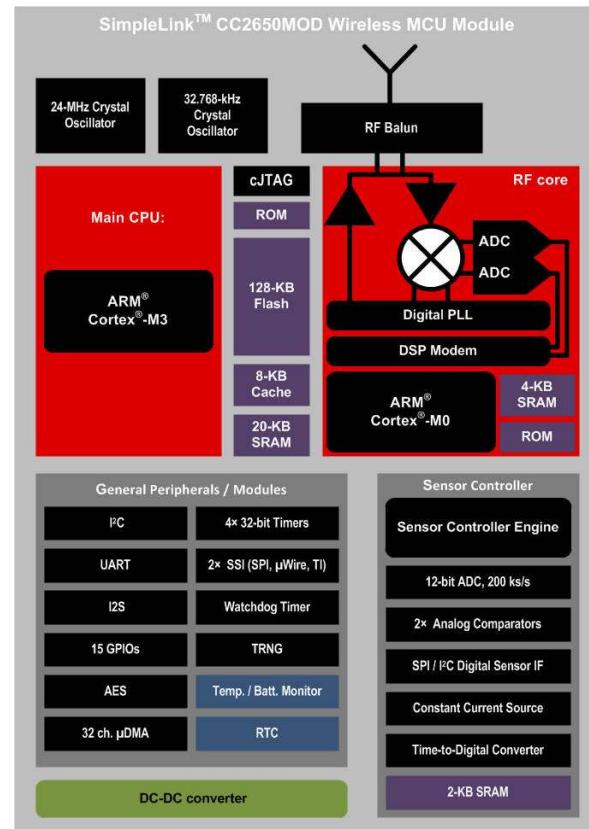
# Implications on design

- A temperature control system using an MPU and MCU
  - Integration costs?



# Some other commercial example

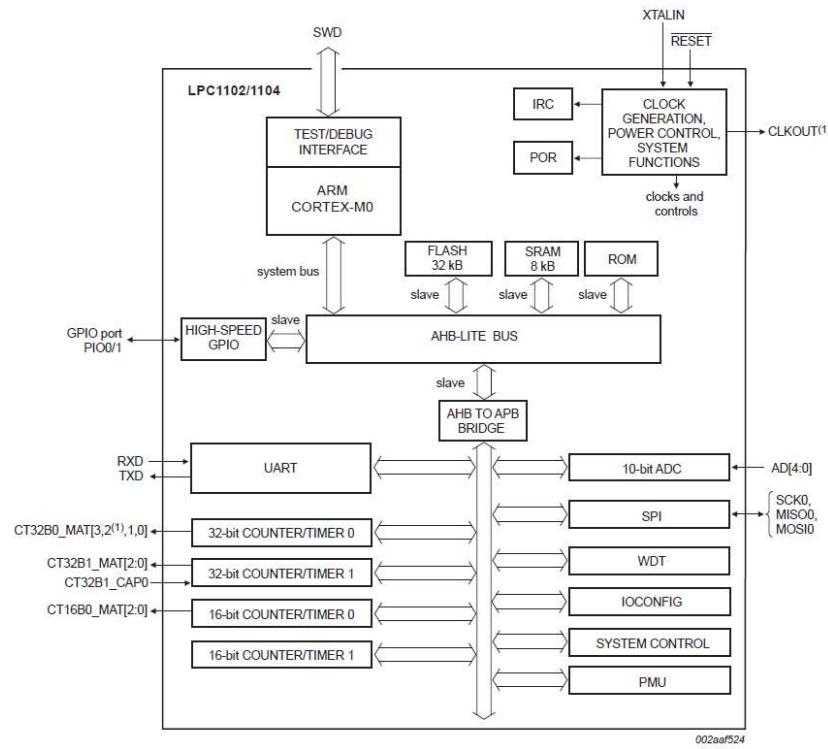
- TI's CC2650
  - "High-end" ultra low-power MCU
  - M3 core (16 bit) (<48MHz)
  - Includes wireless Xceiver
  - 6mA Rx/Tx
  - 60uA/MHz!
    - 2mA @max speed
  - 2.4-GHz RF Xceiver  
Compatible With BLE  
and IEEE 802.15.4



Copyright © 2016, Texas Instruments Incorporated

# Some other commercial example

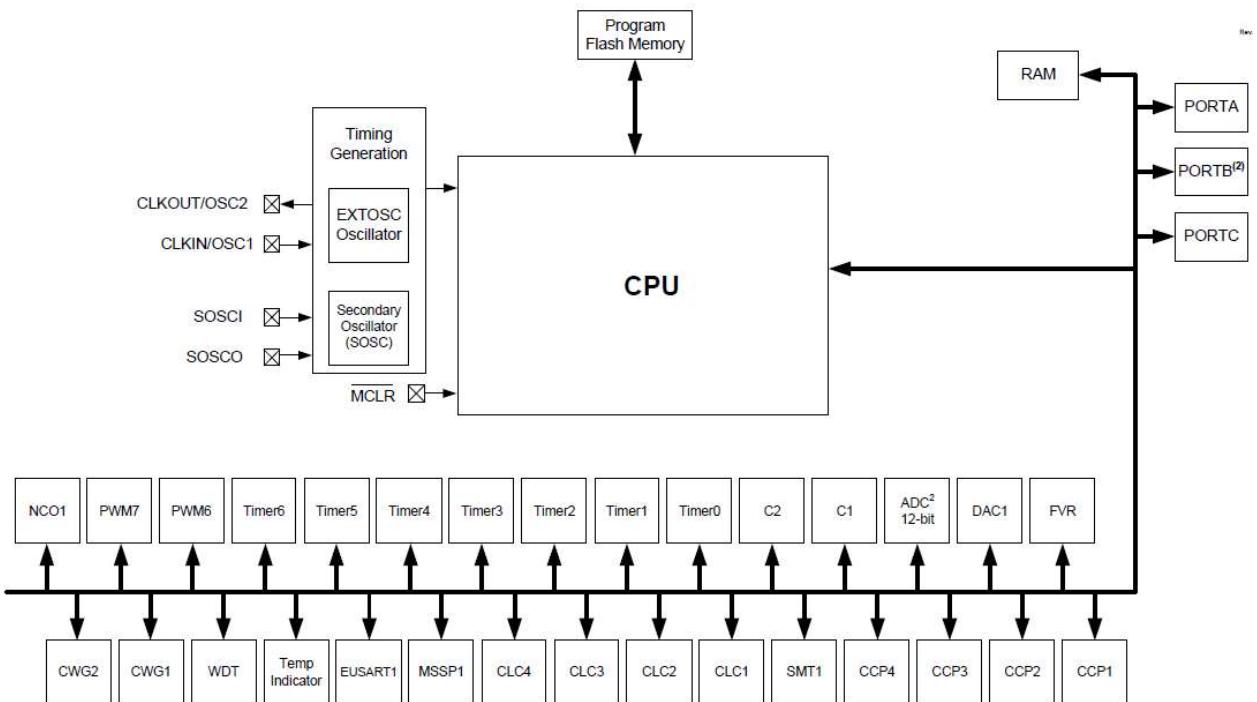
- NXP's LPC1102/1104 (32-bit)
  - M0 core
  - 12MHz-50MHz Clock
  - V<sub>supply</sub> = 3.3V
  - On current = 2-7mA!



# Some other commercial example

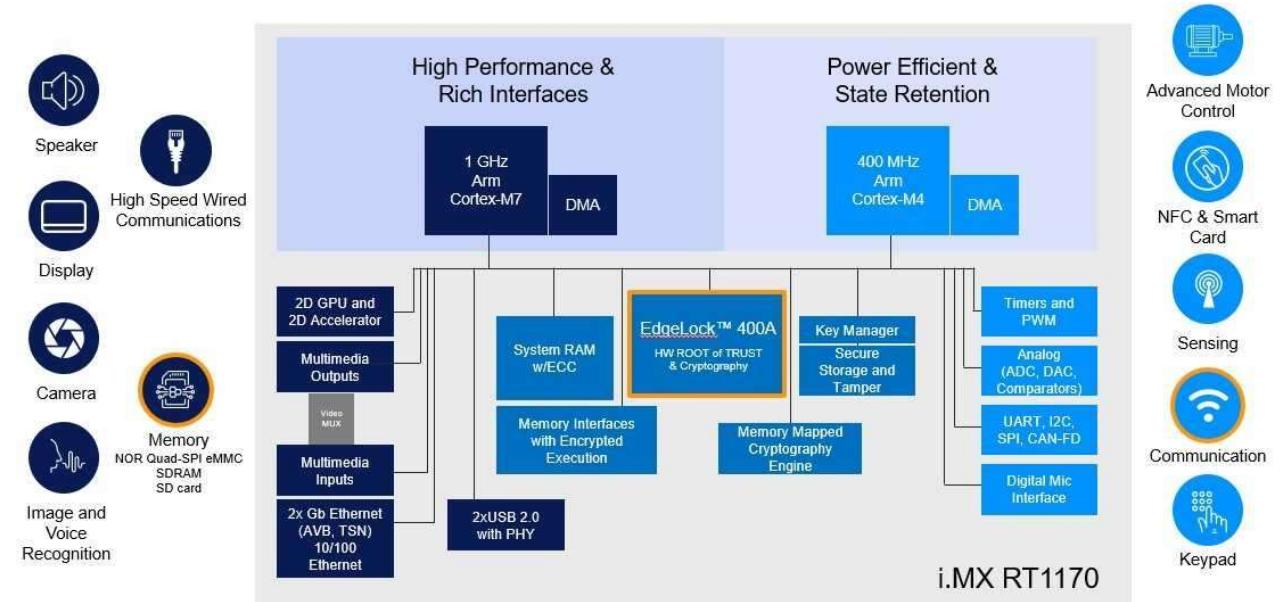
- Microchip's PIC MCU
  - 7KB RAM!
  - 16 bit
  - On-current: 2-4mA @3V

CCP = Capture, Compare & PWM  
CWG= Complementary Waveform Generator  
SMT = Signal Measurement Timer  
CLC = Configurable Logic Cell  
WDT = Watchdog Timer



# Data-analytics is changing the rules

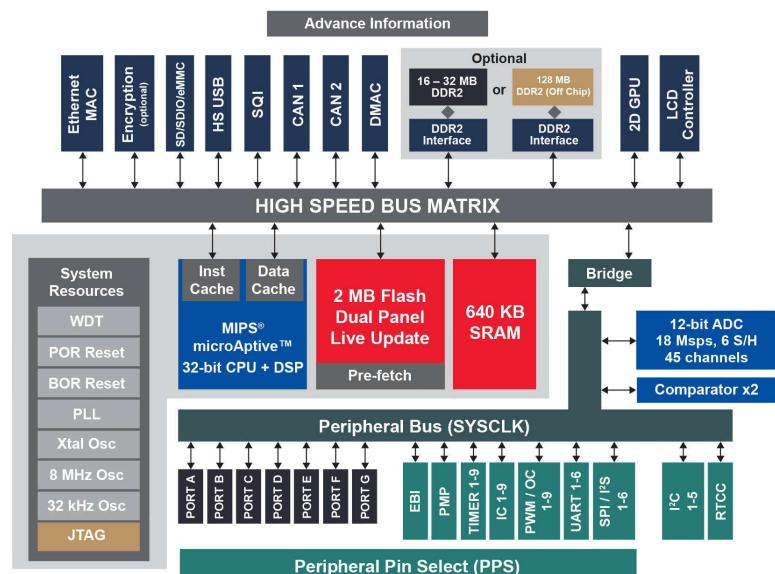
- MCUs meet high performance
- Process data at the source



# MCUs breaks GHz barrier



PIC32MZ DA Family



<b>Processor</b>	Arm® Dual Cortex® -A7 650 MHz L1 32kB I   L1 32kB D 256kB L2 Cache	Arm® Cortex® -M4 209 MHz FPU   MPU
<b>External Memories</b>	DDR3/DDR3L/LPDDR2/LPDDR3 32-bit @ 533 MHz 3x SDMMC	Dual Quad-SPI 16-bit SLC NAND 8-bit ECC
<b>Internal Memories</b>	MCU System RAM 384kB System RAM 256kB	MCU Retention RAM 64kB Back up RAM 4kB OTP fuse 3kb
<b>Connectivity</b>	10/100M or Gigabit Ethernet GMAC 3x USB 2.0 Host/OTG with 2x HS PHY	3D GPU OpenGL ES 2.0 @ 533 MHz MIPI-DSI controller LCD-TFT controller
<b>Security</b>	Camera interface HDMI-CEC 2x CAN FD MDIO slave DFSDM (8 channels/6 filters) 6x SPI / 3x PS 6x I <sup>C</sup> 4x UART + 4x USART 4x SAI SPDIF	TrustZone AES 256, TDES* SHA-256, MD5, HMAC 3x Tamper Pins with 1 active Secure Boot* Secure RAMs Secure Peripherals Secure RTC Analog true RNG 96-bit unique ID
<b>System</b>	5x LDOs Internal and External Oscillators MDMA + 2x DMA Reset and Clock 3x watchdogs Up to 176 GPIOs	5x LDOs Internal and External Oscillators MDMA + 2x DMA Reset and Clock 3x watchdogs Up to 176 GPIOs
<b>Control</b>	2x 16-bit advanced motor control timers 15x 16-bit timers 2x 32-bit timers	2x 16-bit advanced motor control timers 15x 16-bit timers 2x 32-bit timers
<b>Analog</b>	2x 16-bit ADCs 2x 12-bit DACs	2x 16-bit ADCs 2x 12-bit DACs

\*available for STM32MP157C only

# Other issues you must be aware

---

- In general, MCU vs MPU is a clear dichotomy
  - MCUs: run control algorithms
  - MPUs: run computationally intensive algorithms
  - You clearly have in mind what you need
    - MCU selection for your application not a difficult task
- To choose is getting less trivial for modern applications
  - Edge-computing: analytics-on-sensors
    - MCUs with MPU capability: heterogenous HW
  - The complexity of SW to be run is the discriminant factor
    - Not just functionality but also performance and power budget become important
      - Volume of data to process
      - kind of data-processing, e.g. deep-learning at the edge
      - Performance and real-time applications

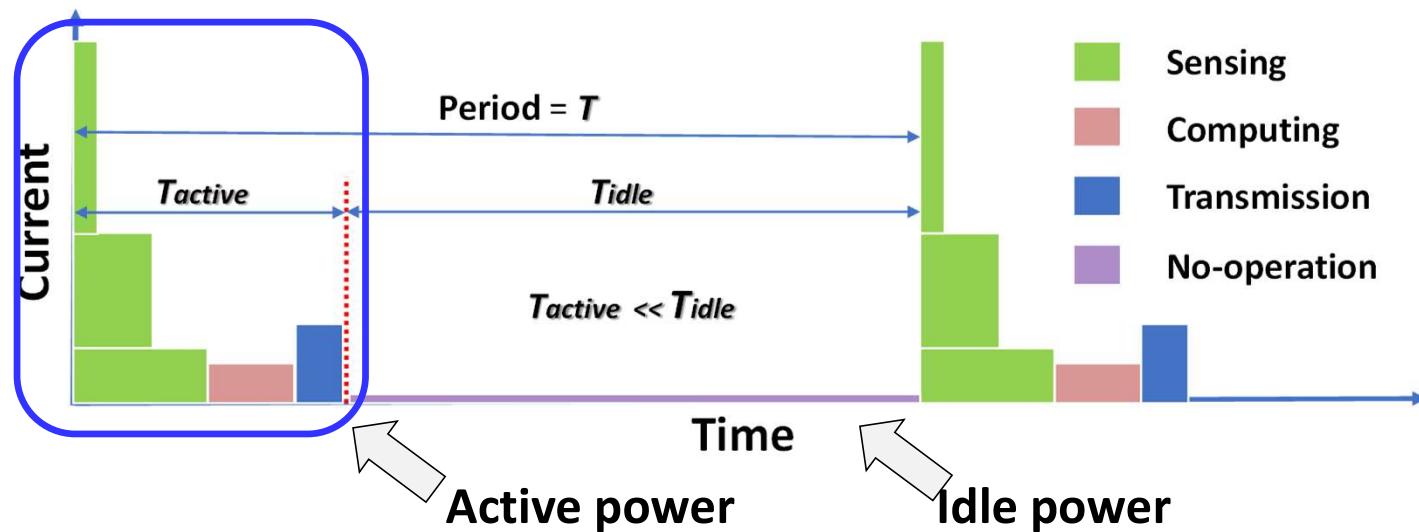
# Other issues you must be aware

---

- For classical (old-style) MCU-based applications, it is also important to consider
  - Programming models/support
    - Forget about fancy IDE with nice GUIs and support for virtually infinite programming languages
      - Simple IDEs are however available...
      - Need to interact with low-level details
  - SW support/IDE/toolchain and programming models and ready-to-use libraries are an important element for the choice of an MCU!
    - Reduce time-to-design
  - Idle power consumption

# Idle power consumption

- For MPUs active power matters
  - Active = when the MPU computes
- For MCUs idle power is at least as important
  - Idle = when the MCU is waiting for events

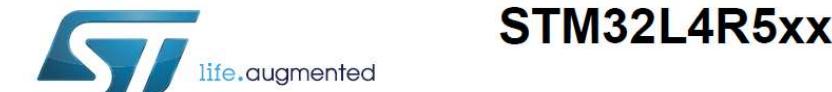


# MCU typical workload

---

- MCU workloads are dominated by idle time...!
  - Small computational tasks (sensing, processing, actuating)
  - Often long duty cycles
- Duty cycle  $D = T_{active}/T$ 
  - No larger than 10-2, often 10-4/10-5
- MCU idle power is a fundamental parameter!!!
  - Notice that there are different types of 'idleness', depending on what is disabled
  - Ex: CC2650
    - Idle/standby/shutdown currents = 0.5mA/1uA/0.15uA
    - The lower, the longer the time to resume to active!  
(14 µs /151 µs / 1015 µs)

# Example: STM32L4R5

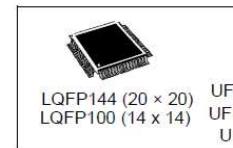


## STM32L4R5xx

Ultra-low-power Arm® Cortex®-M4 32-bit MCUs  
up to 2MB Flash, 640KB SRAM, LCD-TFT & N

### Features

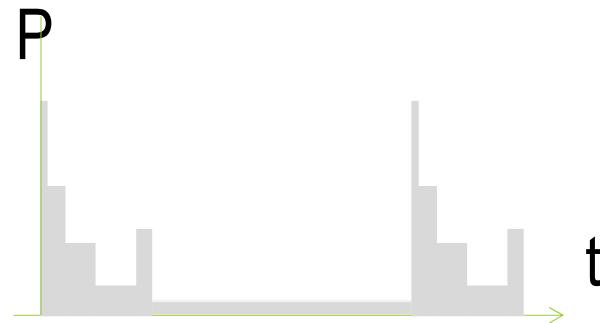
- Ultra-low-power with FlexPowerControl
  - 1.71 V to 3.6 V power supply
  - -40 °C to 85/125 °C temperature range
  - Batch acquisition mode (BAM)
  - 305 nA in VBAT mode: supply for RTC and 32x32-bit backup registers
  - 33 nA Shutdown mode (5 wakeup pins)
  - 125 nA Standby mode (5 wakeup pins)
  - 420 nA Standby mode with RTC
  - 2.8 µA Stop 2 with RTC
  - 110 µA/MHz Run mode (LDO mode)
  - 43 µA/MHz Run mode (@ 3.3 V SMPS mode)
  - 5 µs wakeup from Stop mode



- Internal multi oscillator, aut ±0.25 % acci
- Internal 48 M
- 3 PLLs for sy
- RTC with HW ca
- Up to 24 capacit touchkey, linear
- Advanced graph
  - Chrom-ART / enhanced gr

# Power and Energy: avoid confusion

- Power is an instantaneous metric
  - Energy is cumulative  
 $E(t) = \int P(t) dt$   
(area of power waveform)
- Energy is what you consume from a battery!
  - Depends on power AND time

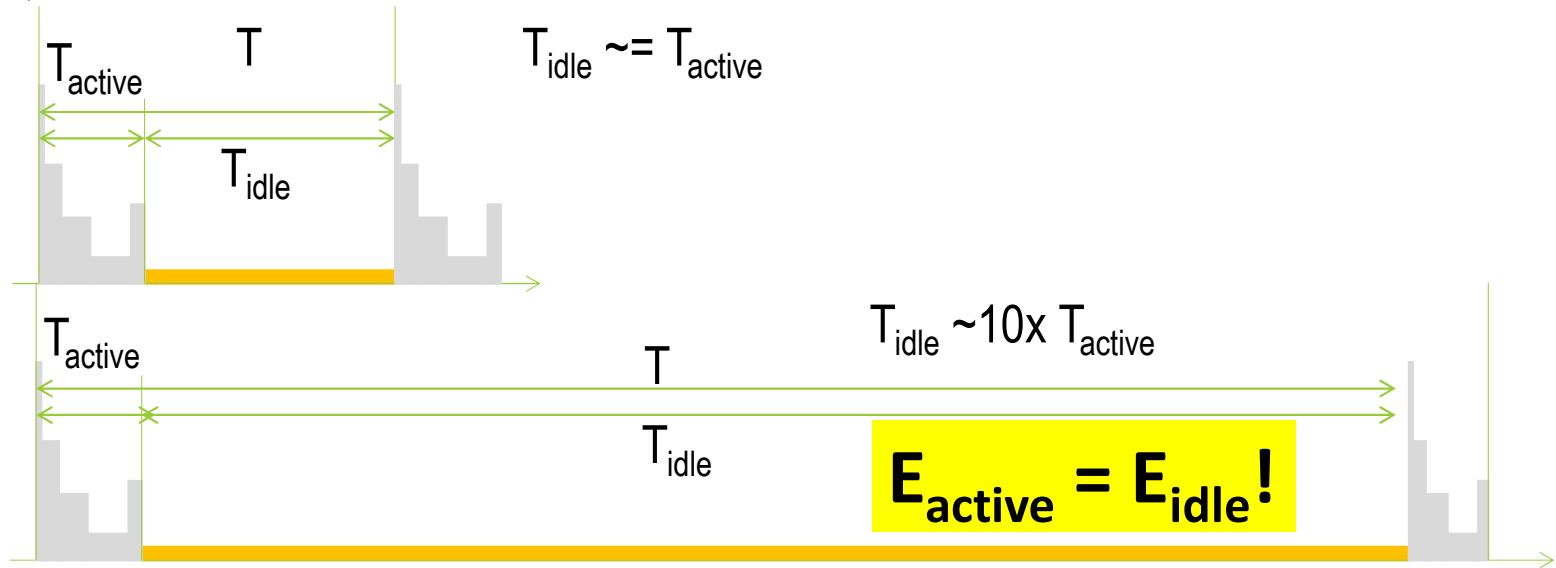


# Importance of duty cycle

- Relevance of static power  $\propto$  the ratio  $T_{\text{idle}}/T_{\text{active}}$
- Example: 1/100 duty cycle (idle=99%)

$$P_{\text{active}} = 10\text{mW}$$

$$P_{\text{idle}} = 10\mu\text{W} \quad (1/100)$$



Radio communication

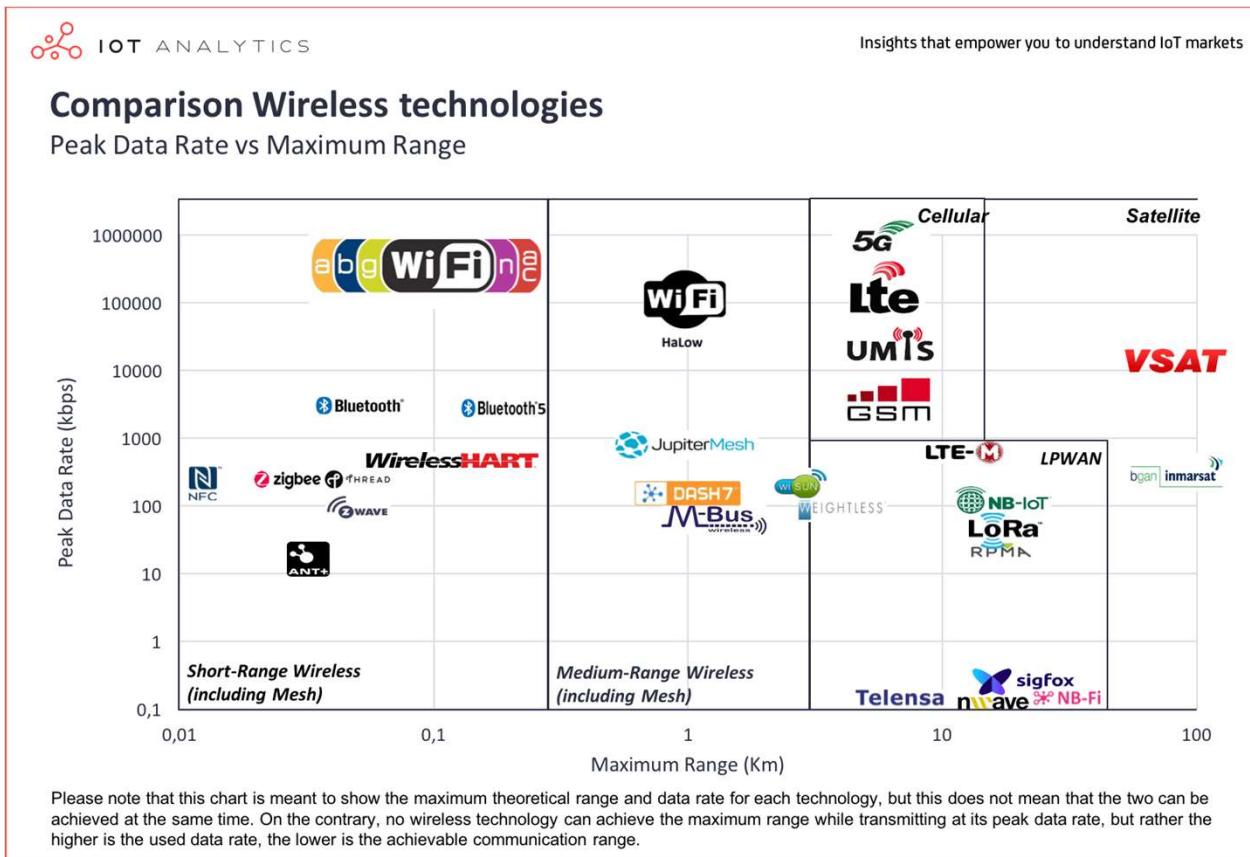
---

# Wireless interfaces

---

- Can be seen as actuators: Antenna as a transducer
- Standardization helps!
  - Wireless interfaces allow
- Many processing cores include radios and wireless I/F
  - But we need to understand which is the wireless communication needed by our application
- FOMs
  - Bandwidth/bitrate: correlated to data volume
  - Power consumption
  - Range
  - Cost
- Other important characteristics
  - Architecture: Point-to-point vs. broadcast, # of supported nodes

# Options available



# Comparison chart

	Range	Standard	Max nodes	Tx Speed	Power/Energy	Price	Frequency
<b>ANT</b>	< 30m	Propriet.	$\infty$ ( $2^{32}$ )	1Mb/s	<b>~100 uW [700 nJ/bit]</b>	5\$	2.4 GHz
<b>Bluetooth (LE)</b>	10m (100m)	Open	7+1	O(1 Mb/s)	<b>10mW, &lt;15mA [70nJ/bit]</b>	5\$	2.4 GHz
<b>NFC</b>	< 30cm	Open (multiple)	1 (M/S)	106 to 424 kbps	<b>&lt;15mA</b>	10c	13.56MHz
<b>UWB</b>	< 10m	Open	7+1	110/480 Mb/s	<b>~200 uW</b>	\$\$\$	3.1 to 6 GHz
<b>Wifi</b>	< 100m	Open	32	54 Mb/s	<b>250 mW @ 1Mbps [5.25nJ/bit]</b>	>10\$	2.4 and 5 GHz
<b>ZigBee</b>	<100m	Open	$\infty$ ( $2^{64}$ )	250 Kbps	<b>O(10) mW [360 nJ/bit]</b>	<5\$	2.4 GHz

# Long-range interfaces (distributed systems)

- An important segment for IoT are long-range radios (LPWANs)

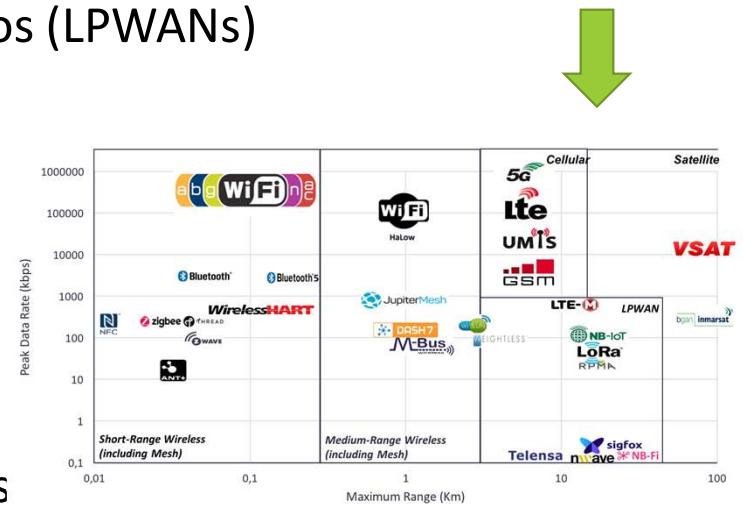
- **Tens of Km range, low-power, low-bitrate**

- Key protocols
  - SigFox, LoRaWAN, NB-IoT

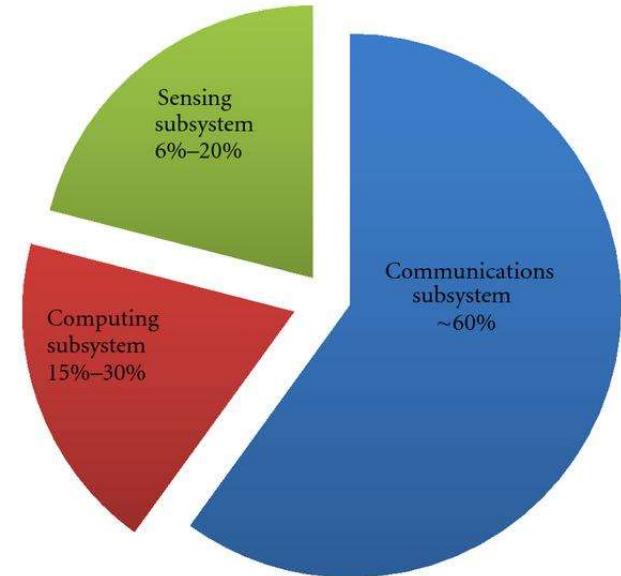
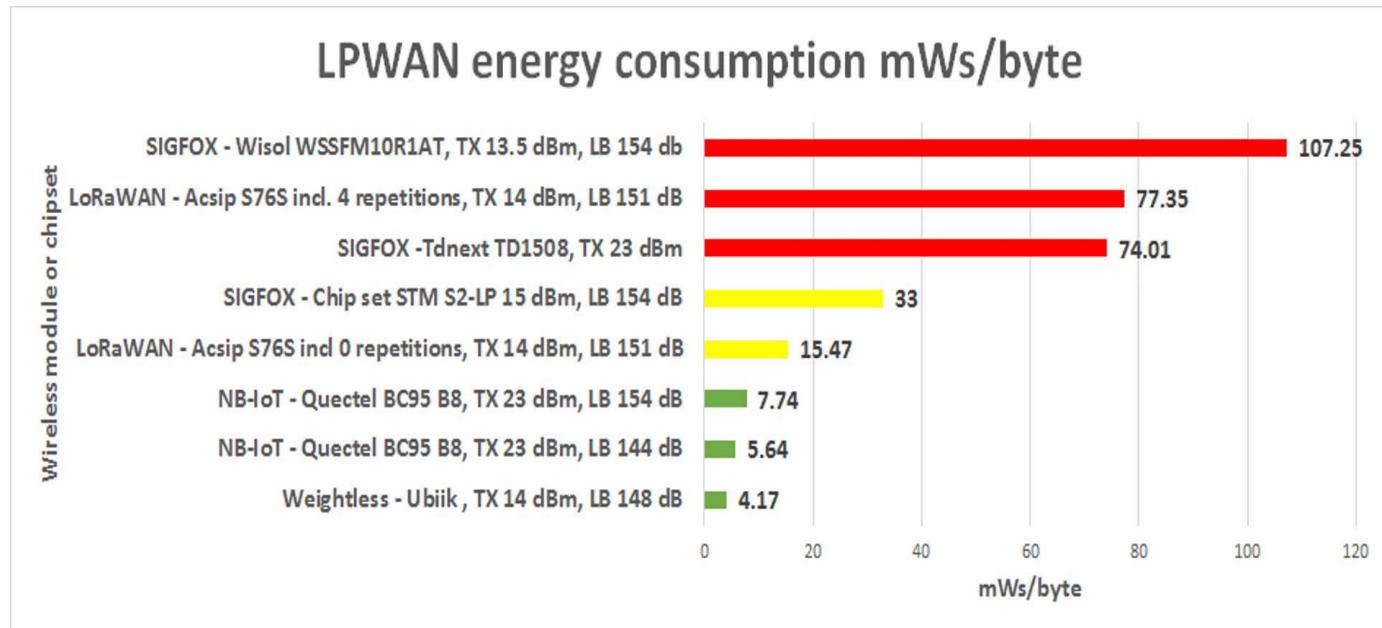
- Operate in sub-GHz bands with very narrow bands

- At those frequencies, signals penetrate obstacles and travel long distances while drawing relatively little power

- Bitrate from **100 bps** (SigFox) to **200 Kbps** (Nb-IoT)

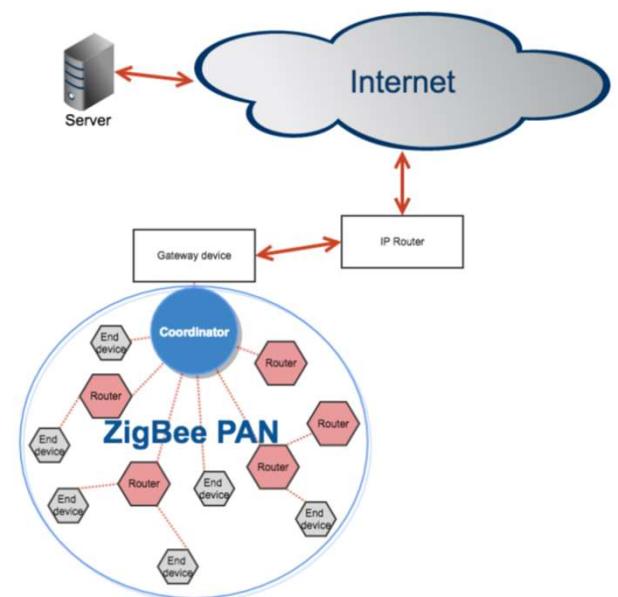
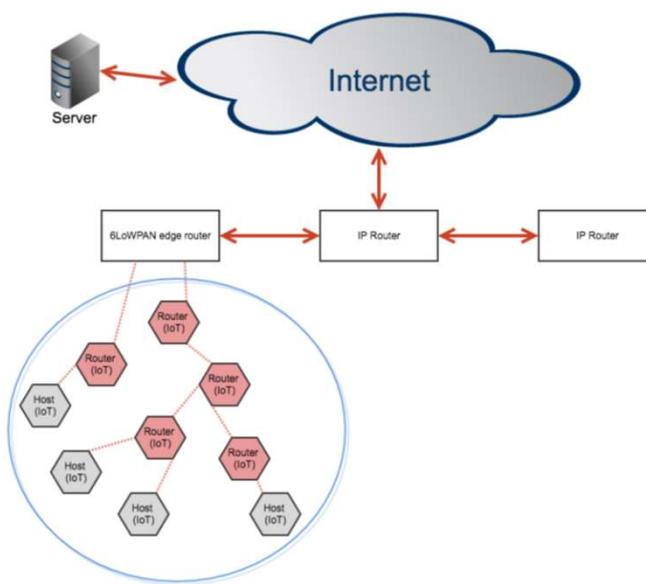


# Energy consumption

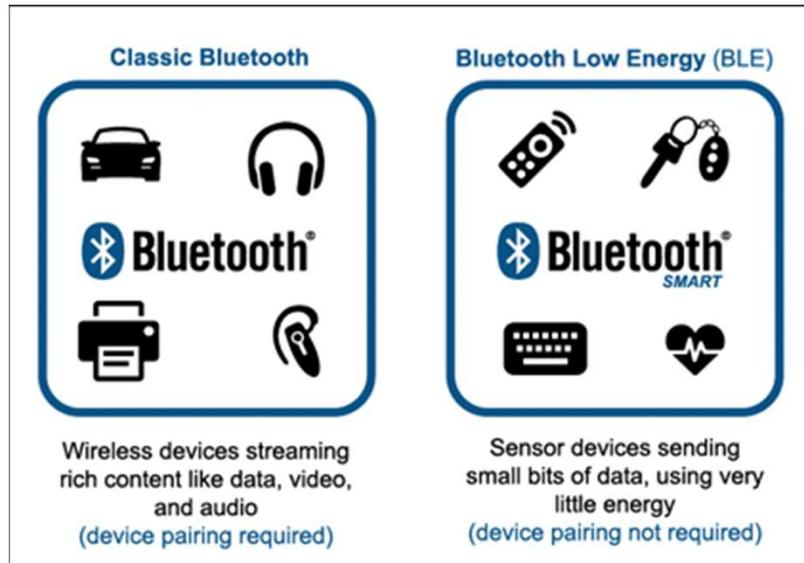


# Mid-range: Wireless Local Area Networks

- Need a coordinator device that initiates network formation.
- More reliable and secure
- Less power
- Slower



# Near, p2p communication



Energy

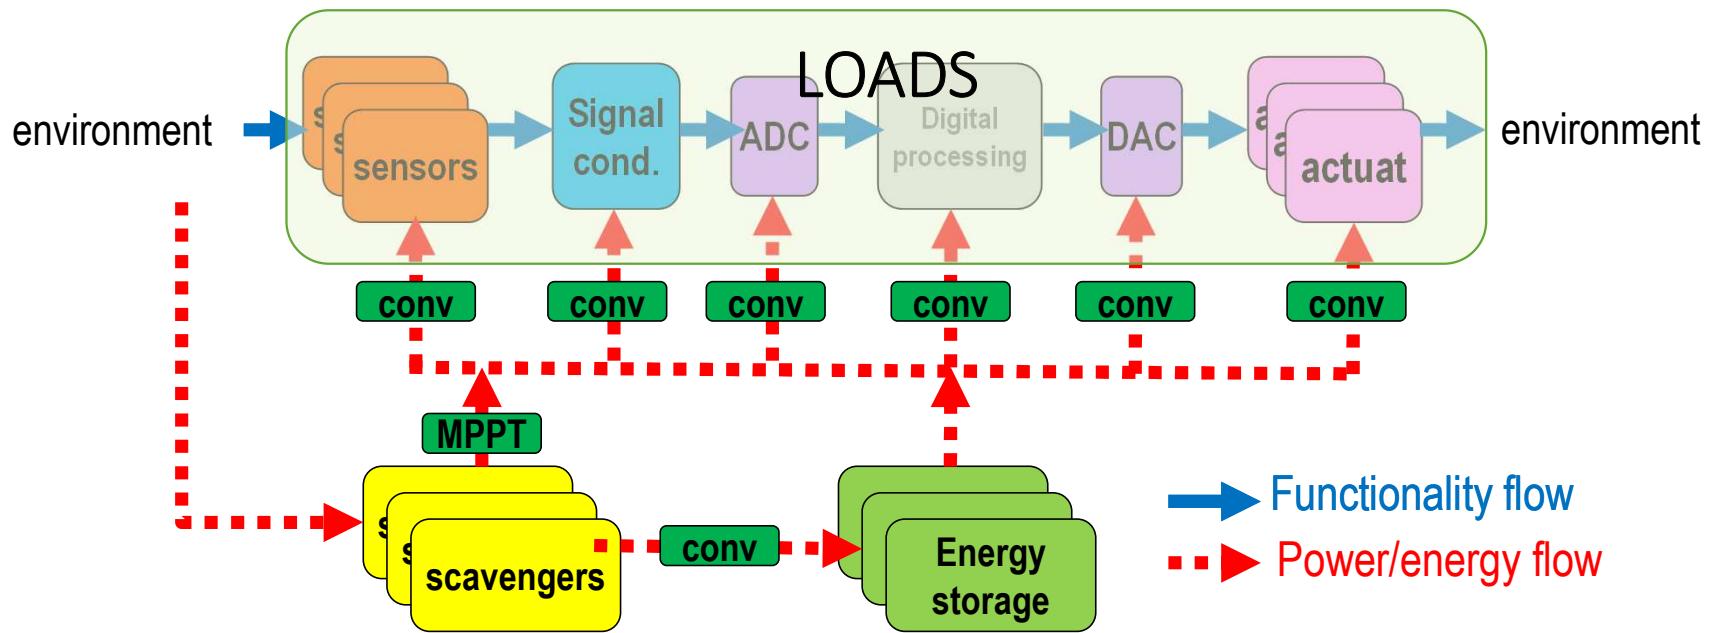
---

# Why bother

---

- Many drivers for optimizing/controlling power consumption
  - Technological
    - Difficult to manufacture devices with high power consumption
    - Batteries do not catch up
  - Market
    - How many devices you normally use are wired...?
  - Economical
    - Costly to manufacture devices with high power consumption
  - Environmental
    - Green computing....
    - Energy is a limited resource
    - Total Energy of Milky Way Galaxy: 1059 J !

# Functional and energy/power flows



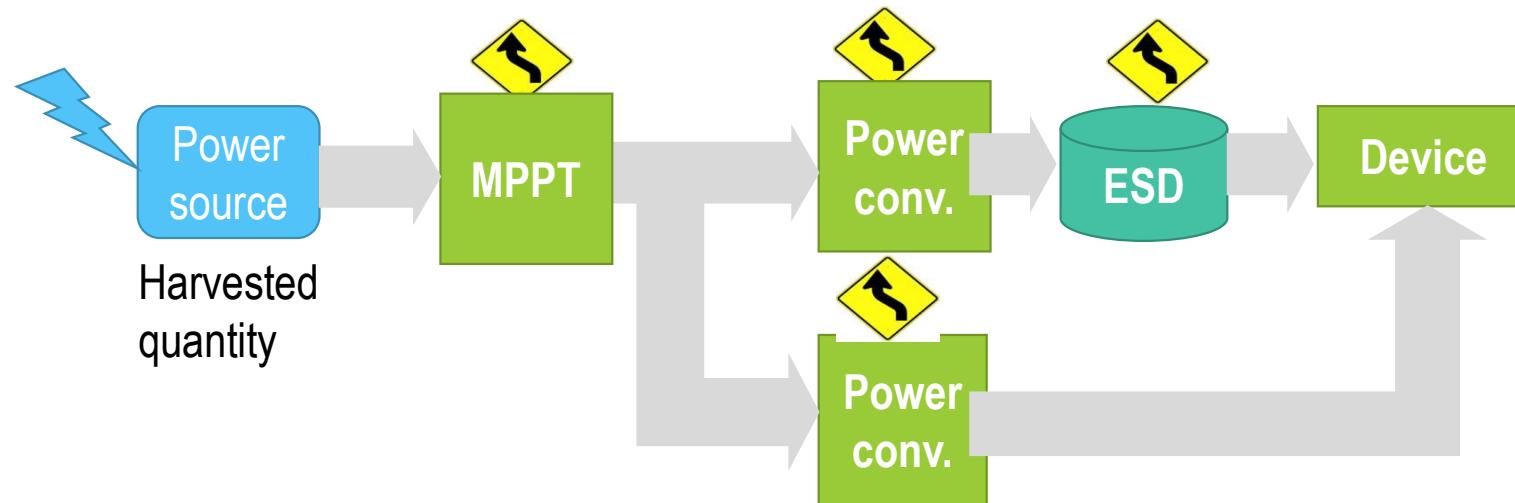
# Energy is a resource to manage

---

- Functionality (@ low power) as commodity!
- Digital and non-digital functional parts (MCU, accelerators, memories, RF, I/Fs) are assumed to be power-optimized
  - Designer picks components, protocols, etc. based on needs and based on their power scores
  - System designer does not design the components but only decides the way they are used (SW)
- Important to reduce energy consumed, but (at least as important as) to manage how energy is:
  - Generated through various types of energy harvesters (power sources)
  - Distributed/converted to the various units
  - Stored in appropriate energy storage devices (ESDs) to guarantee sustained service

# Energy management systems

- Accurate power/energy planning is non-trivial
  - Many points of losses
  - Sometimes contrasting...
  - Simple models can do the job of a quick assessment to avoid gross design mistakes!



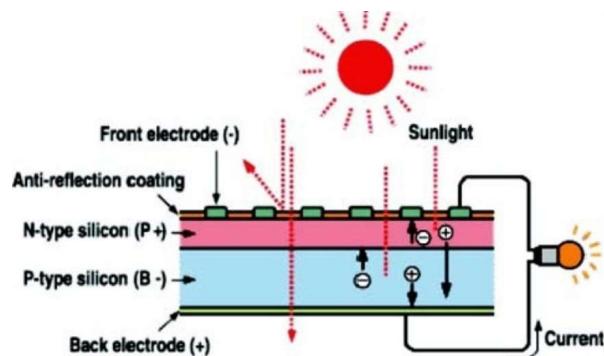
# Energy harvesting

---

- = extracting energy from the environment
- We need electrical energy, so any harvester embeds some kind of transducer
- A large variety of physical effects can be exploited
  - Piezoelectric, photoelectric, triboelectric, thermoelectric,...
- As users, we are more interested in the “sources” from which we can extract energy...
- (solar) radiation
  - Vibration
  - temperature
  - friction
  - Ambient radio frequency/noise
  - Ambient airflow

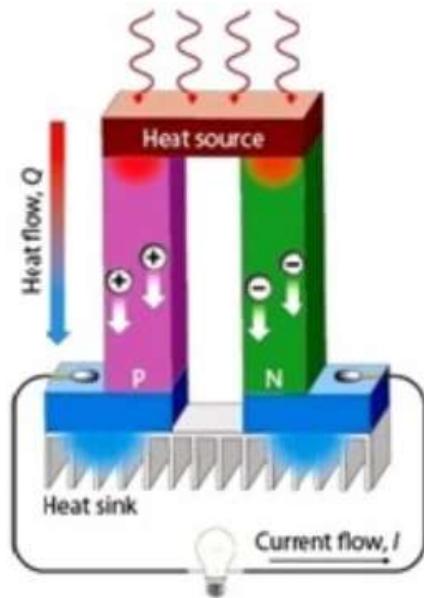
# Harvesting from solar radiation

- Photovoltaic effect
  - Photovoltaic materials when exposed to light (photons) generate current
  - i.e., photons force electrons in the conduction band
- Silicon (P-n junctions) is one such material
  - A PV-cell is basically a diode and obey the basic diode rules
  - But reference current is determined by solar irradiance



# Harvesting from temperature

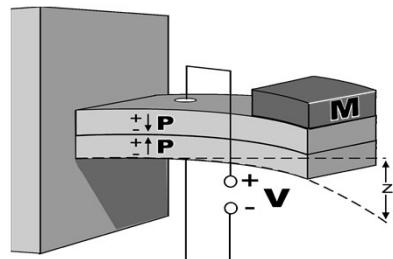
- Thermoelectric effect a.k.a. Seebeck effect
  - A potential is generated when the junctions of two dissimilar metals experience a temperature difference
  - Potential proportional to  $\Delta T$
- Inverse of Peltier effect (thermocouple)



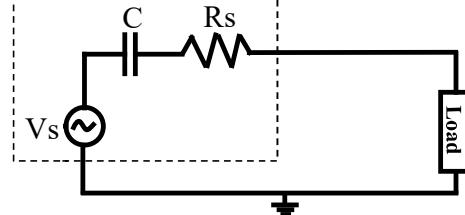
# Harvesting from vibrations

## Piezoelectric

Strain in piezoelectric material causes a charge separation (voltage across capacitor)

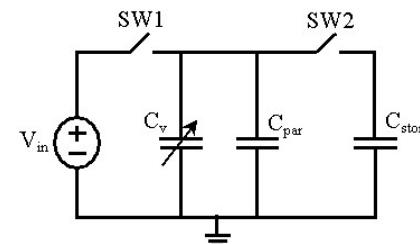
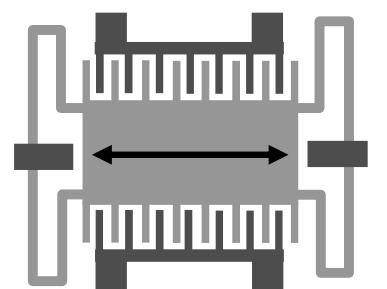


## Piezoelectric generator



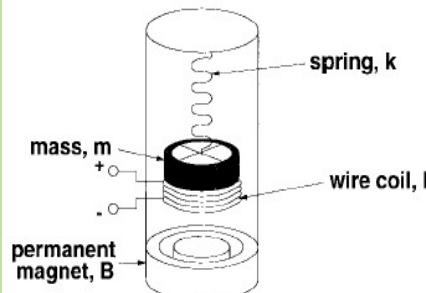
## Capacitive

Change in capacitance causes either voltage or charge increase.



## Inductive

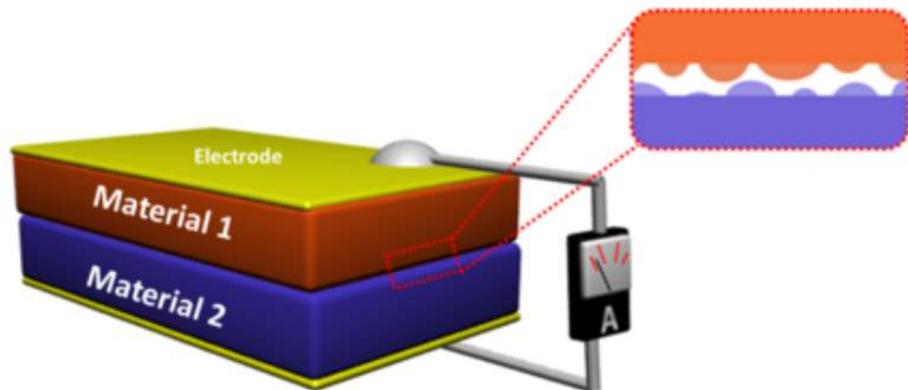
Coil moves through magnetic field causing current in the wire.



Amirtharajah et. al., 1998

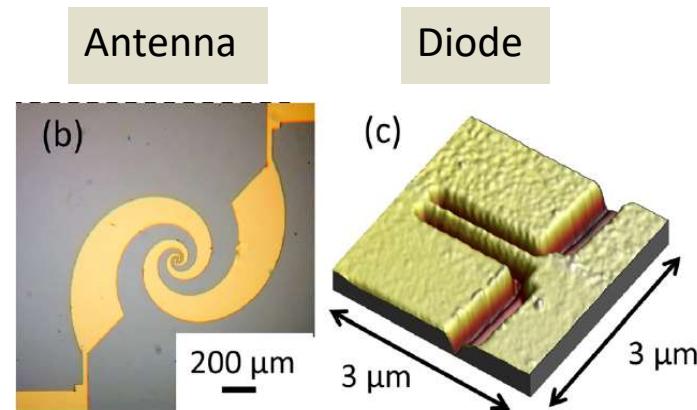
# Harvesting from friction

- Triboelectric effect
  - Friction between materials can generate a potential
  - A particular instance of electrostatics...



# Harvesting EM noise

- EM energy easy to convert into just DC electrical power
  - Typically, an antenna + rectifier will do the job!
- **Rectenna** (rectifying antenna)
  - Micro-scale rectennas are of interest!



# Issues with energy harvesting

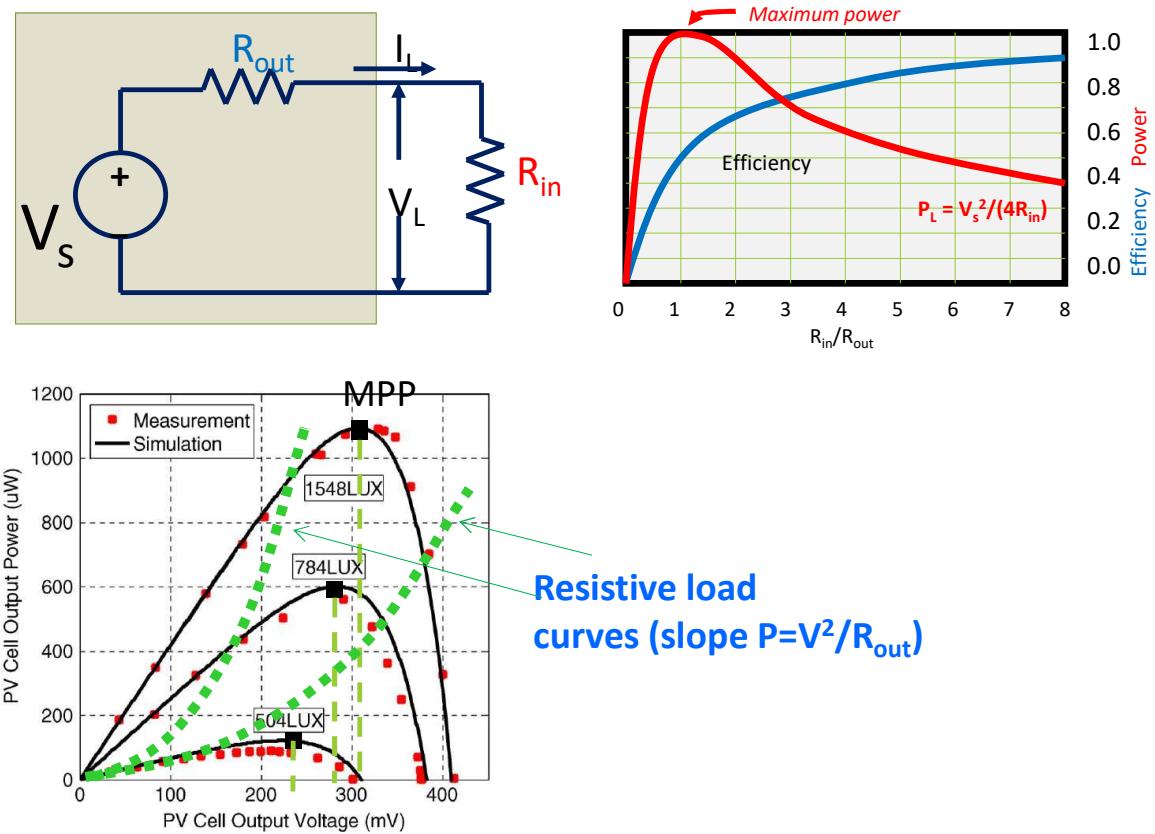
1. Generated power depends on the value of some environmental quantity
  - Q: How do we guarantee as much extracted power as possible under different environmental conditions?
  - A: The “maximum power extraction” problem
  - Q: How can we sustain computing under intermittent energy?
  - A: Intermittent computers
2. Extracted power is almost invariably very low...
  - Q: How can we derive acceptable power levels for typical computing blocks?
  - A: Appropriate power conversion, Scavenger “packing”
3. Extracted power might not be needed “right now”
  - Q: How can we accumulate excess energy?
  - A: Energy storage design & management

Harvesting method	Power density
Solar cells (outdoor at noon)	15 mW/cm <sup>2</sup>
Solar cells (illuminated office)	100 µW/cm <sup>2</sup>
Piezoelectric (shoe inserts)	330 µW/cm <sup>3</sup>
Thermoelectric (10 °C gradient)	40 µW/cm <sup>3</sup>
Acoustic noise (100 dB)	1 µW/cm <sup>3</sup>

A typical battery cell has 100-1000 mW/cm<sup>3</sup> !!!

# Extracting power

- Harvester has low efficiency if storage device is not impedance-matched to source !!!
  - Source impedance is variable
  - Maximum power point tracking (MPPT) problem
- Need proper circuits
  - Tradeoff between energy extracted and consumed by MPPT circuit



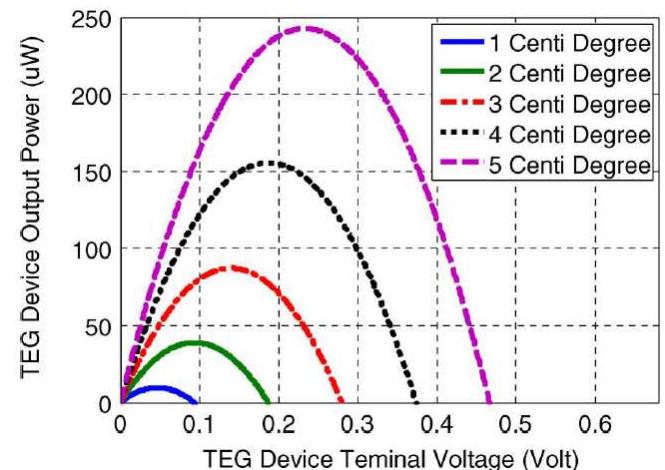
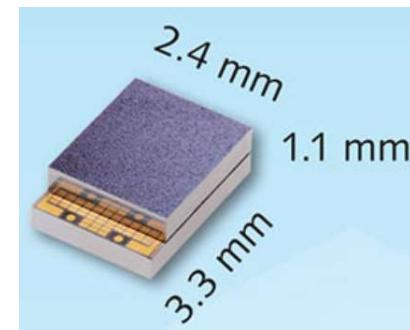
# Tracking the MPP

---

- Done by ad-hoc HW call MPP tracker (MPPT)
- Basically a switching DC-DC converter that tries to match  $R_{in}$  and  $R_{out}$  by adapting the working point
- Double source of inefficiency !
  - Converter efficiency
    - $V_{mpp}$  is not the level you need
  - Failure to perfectly track MPPT

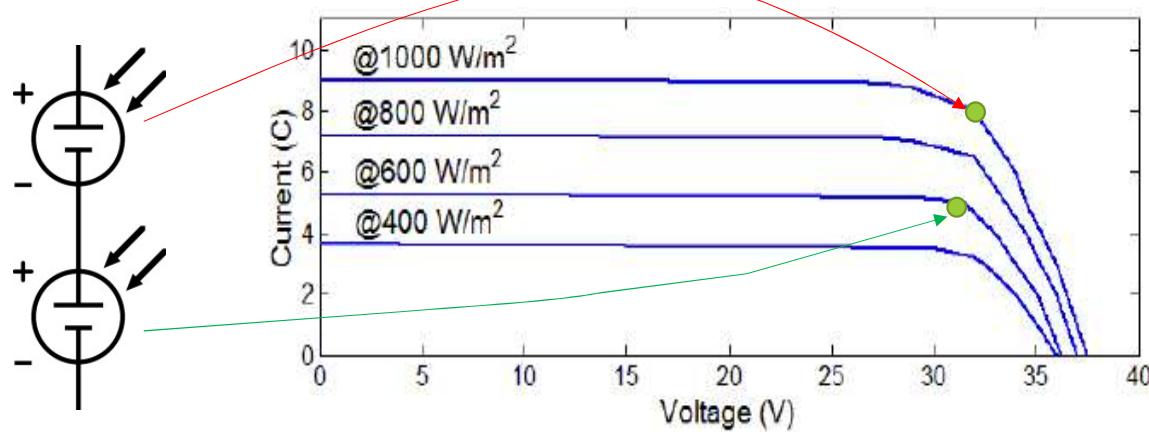
# Extracting from ultra-low power sources

- Example:
  - TEG device Micropelt MPG-D751
  - MPPT for DT=5oC is only 50  $\mu\text{W}!!!$  (@0.25V!)
- How to make such a low power level usable for a typical load (e.g., 10mA @ 1V = 1mW?)
  - Use appropriate converters
- If possible, combine power sources into “modules”
  - scale up both voltage and current
    - Series connection → increase voltage
    - Parallel connection → increase current



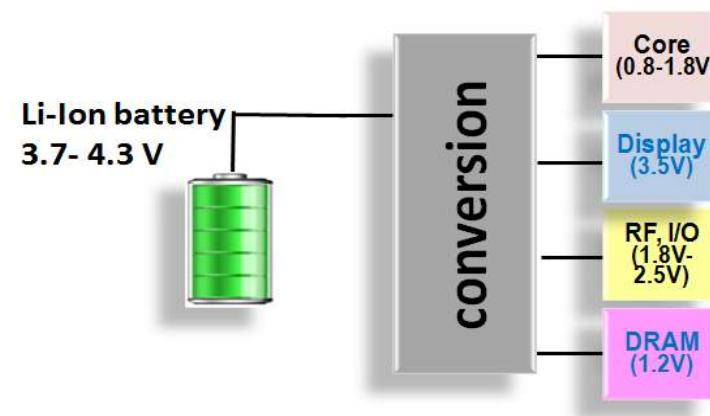
# Packing power sources

- Obviously, space constraints should allow this
- Issues with unmatched devices !!!!
  - i.e., exposed to different environmental conditions
  - Example:



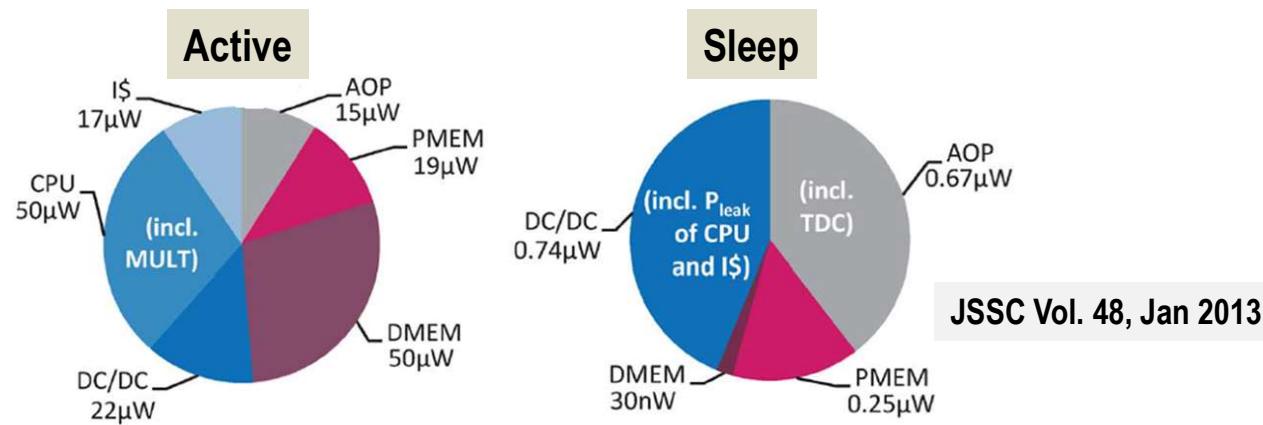
# Power Converters

- a.k.a. Voltage regulators
  - The equivalent of transformers in DC!
- INPUT : an unregulated DC voltage Vin.
- OUTPUT: a regulated output voltage Vout, having a magnitude (and possibly polarity) that differs from Vin
- Conversion is ubiquitous!!!
  - Need it anytime you need to interface two power levels



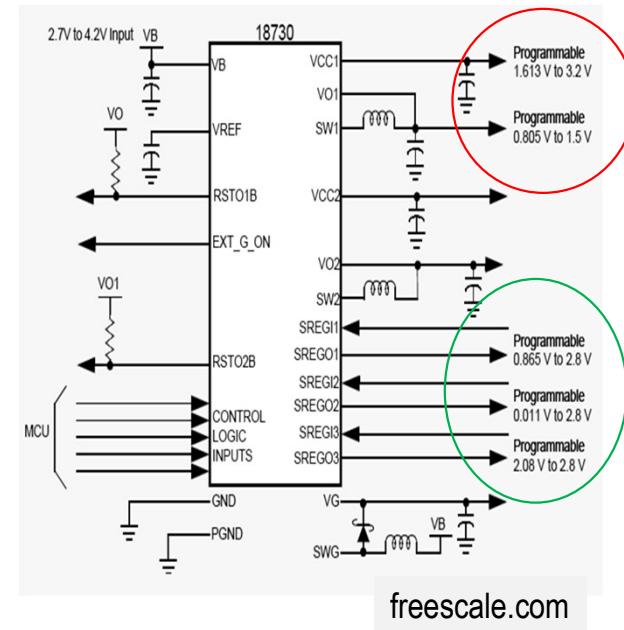
# Example figures

- SleepWalker node from IMEC
  - Actually a ULP µcontroller (7 $\mu$ W/MHz) including a DC/DC converter
- DC/DC losses take from about 12% to 40%
  - Small but even more significant in sleep state...
  - And this is just one converter...



# Example figure

- Freescale's MPC18730 Multi Channel DC/DC Converter
  - 2 DC/DC Converters + 3 LDO (low-dropout) Regulators
- Example:
  - input from Li-Ion Battery 4.0V
  - DC/DC
    - $V_{out} = 1.2V @ 50mA = 60mW$
    - $V_{out} = 2.5V @ 100 mA = 250mW$
  - LDOs
    - $V_{out}=1.8V @ 50mA = 90mW \times 3$
- Power loss in conversion
  - DC/DC = 85% efficiency
  - LDO =  $V_{out}/V_{in} = 45\%$  efficiency



# Example figure

- Freescale's MPC18730 Multi Channel DC/DC Converter
  - 2 DC/DC Converters + 3 LDO Regulators

- Example:

- in

Total loss:

- D

DC/DC1 = 44mW

- L

DC/DC2 = 10mW

- LD

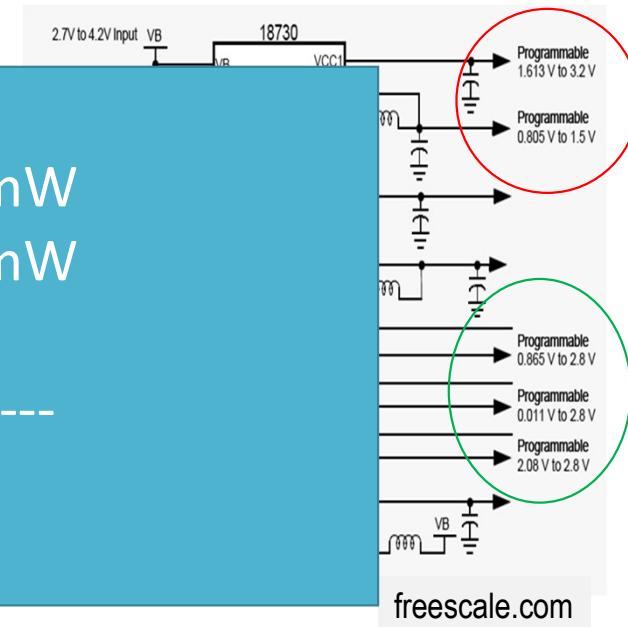
LDOs =  $110\text{mW} \times 3 = 330\text{mW}$

- Power

**Total losses 384mW**

- D

**(66% of consumed power !!!)**

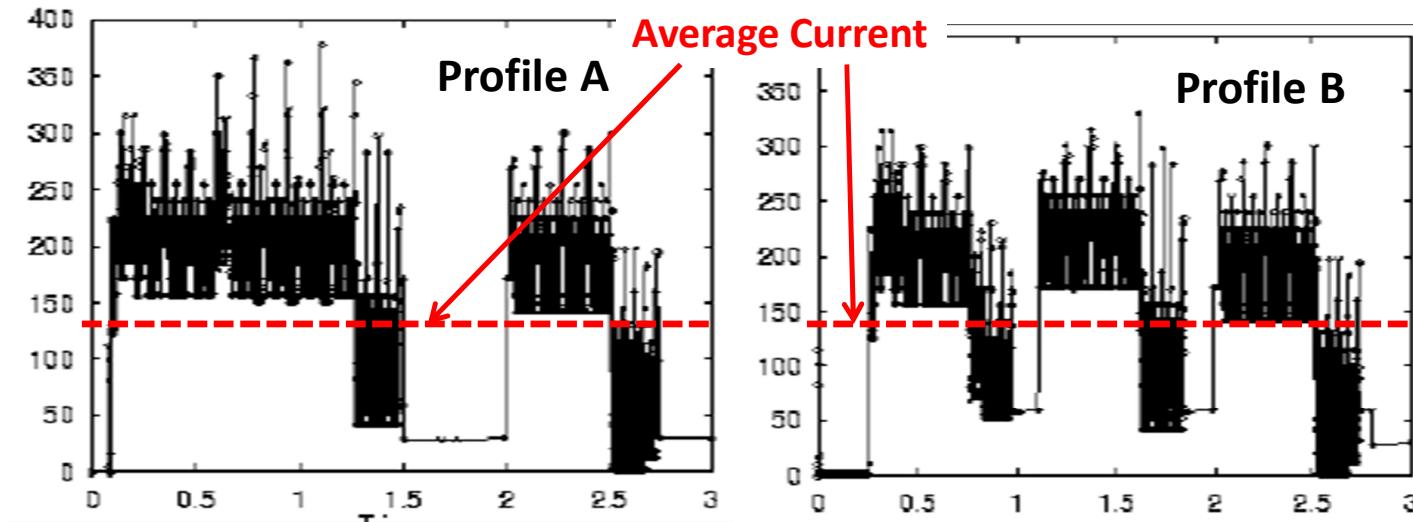


# Energy Storage Devices (ESDs)

---

- Most typical embodiment of ESD is a battery
- Many possible alternatives (even at mscale)...
  - Ultra-capacitors
  - Fuel cells
  - Flywheels
  - ...

# Battery discharge path



[Source: Intel]

Profile	Average Current(mA)	Battery Life (ms)	Specific Energy(Wh/Kg)
A	123.8	357053	15.12
B	124.2	536484	18.58

# Super- or Ultra-capacitors

- Supercapacitors (also known as ultracapacitor or double-layer capacitor) differs from a regular capacitor in that it has a very high capacitance
- But it is indeed a capacitor....
  - It stores energy by means of a static charge (as opposed to an electrochemical reaction)
  - Applying a voltage differential on the two plates will charge the capacitor



**120F, 2.5V supercap –  
20US\$**

# Supercaps vs batteries: hybrid ESD

Figure of merit	Supercaps	Li-Ion (generic)
<b>Charge time</b>	Seconds	Hours
<b>Cycle life</b>	$10^6$	<1000
<b>Cell voltage</b>	2.3 – 2.75 V	3.6 – 4.3 V
<b>Specific Energy [Wh/kg]</b>	5-10	100-200
<b>Specific Energy [W/kg]</b>	Up to 10,000	1000 to 3000
<b>Cost per Wh</b>	\$10-20	\$1-2
<b>Charge temperature</b>	-40 to 65 °C	-0 to 45 °C
<b>Discharge temperature</b>	-40 to 65 °C	-20 to 60 °C